



# CS 415

# Operating Systems

# System Calls

Prof. Jieyang Chen

Department of Computer and Information Science

Fall 2025



UNIVERSITY OF OREGON



UNIVERSITY  
OF OREGON

# *Logistics*

- Project 1 posted
- Labs
  - Intended to support projects
  - Lab 2 posted on Monday
- Start reading OSC Chapter 3
- Make use of Canvas discussion
- Get familiar with the Linux man pages  
<https://man7.org/linux/man-pages/index.html>

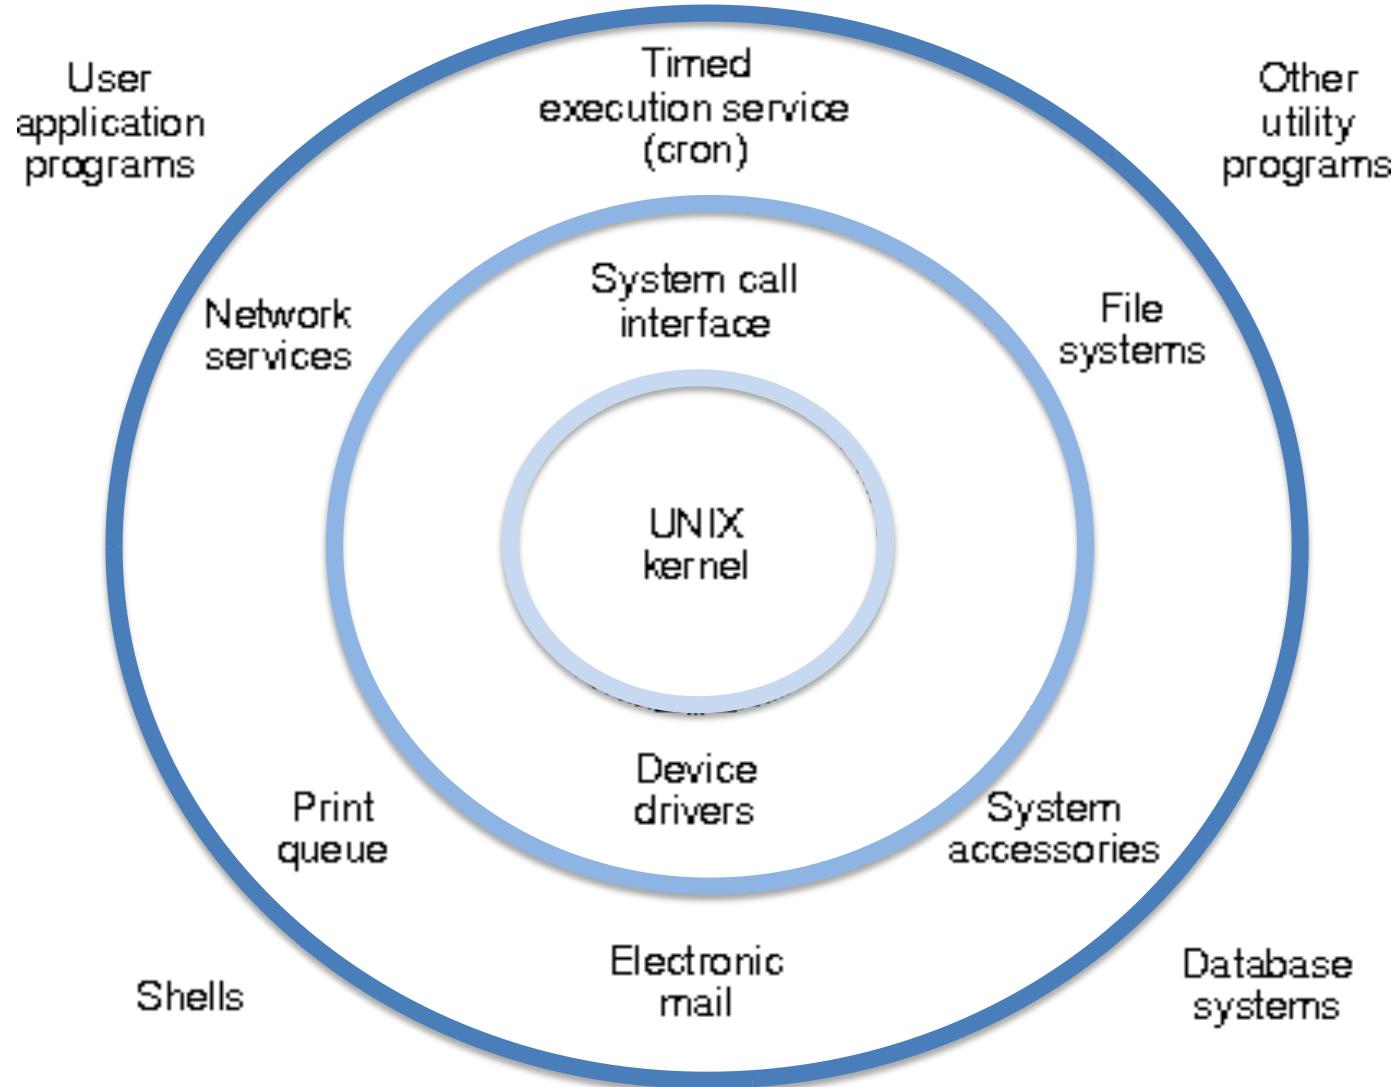
# *Outline*

- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation

# *Objectives*

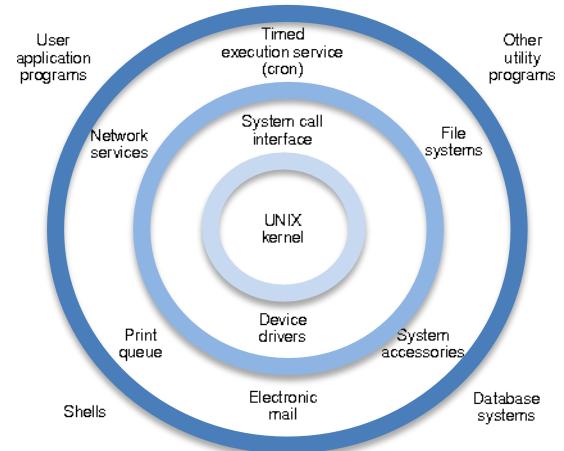
- To describe the services an operating system provides to users, processes, and other systems
- To explain how operating systems handle system calls

# *Operating System Layers*



# *System Layers*

- Application
- Libraries (in application process)
- System Services
- OS API
- Operating system kernel
- Hardware



# *Applications to Libraries*

- Application programming interface (API)
- Libraries
  - Example: *libc*
- Library routines
  - Example: *printf()* of *stdio.h*
- All within the process's address space
  - Statically linked
    - ◆ libraries are included as part of the application code
    - ◆ calls are resolved at compile time
  - Dynamically linked
    - ◆ libraries are loaded by the OS at execution time as needed
    - ◆ jump tables and pointers are resolved dynamically by linker

# *Libraries to System Routines*

- System call interface
  - UNIX man pages, section 2
  - Examples
    - ◆ *open()*, *read()*, *write()* – defined in *unistd.h*
  - Call these via libraries? *fopen()* vs. *open()*
  - See links in schedule
- Special files
  - Drivers
  - */proc*
  - *sysfs*

# *System to Hardware*

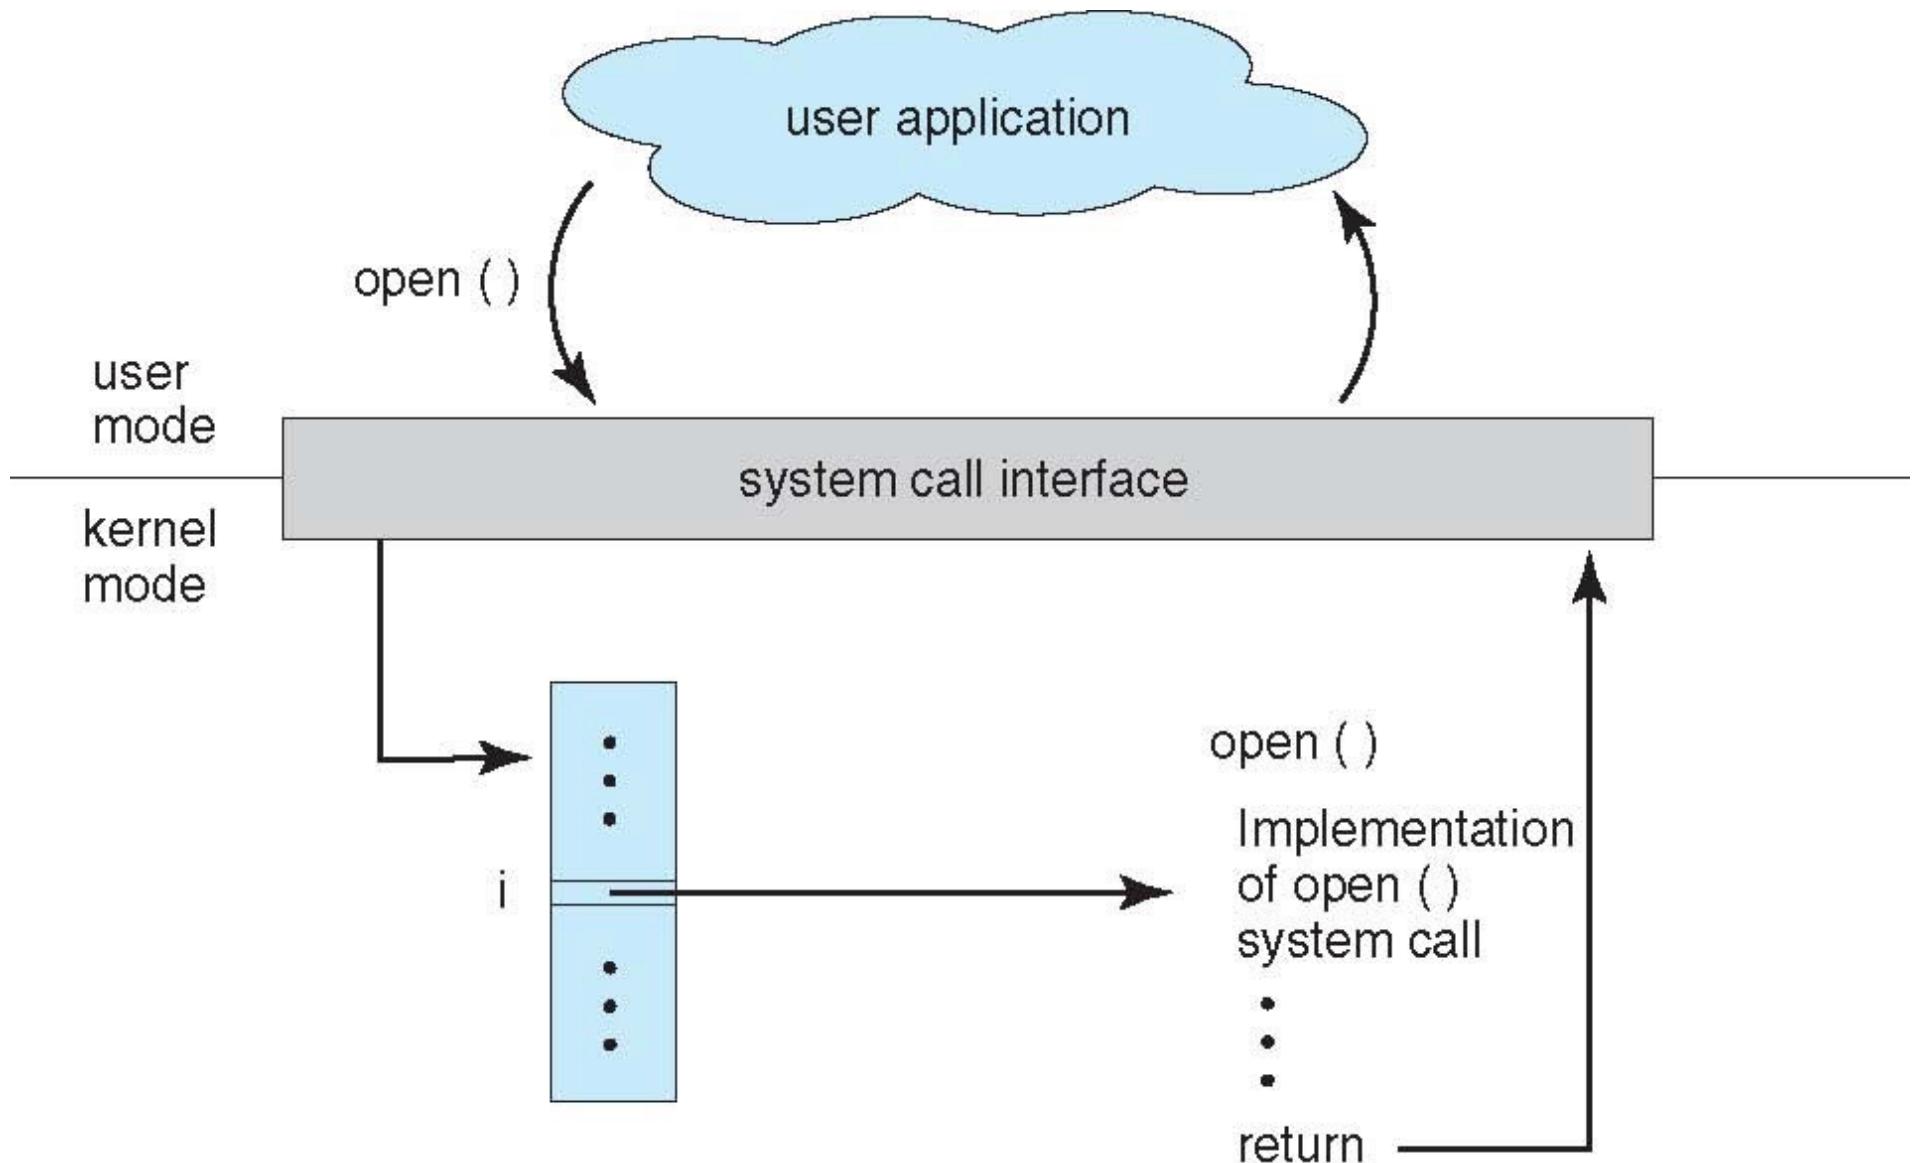
- Software-hardware interface
- OS kernel functions
  - Concepts               Managers (hardware)
  - Files                   File systems (drivers and devices)
  - Address space         Virtual memory (memory)
  - Programs               Process model (CPU, ISA)
- OS provides abstractions of devices and hardware objects
  - These abstractions are represented in software running in the OS and data structures that it maintains

# *Systems Calls*

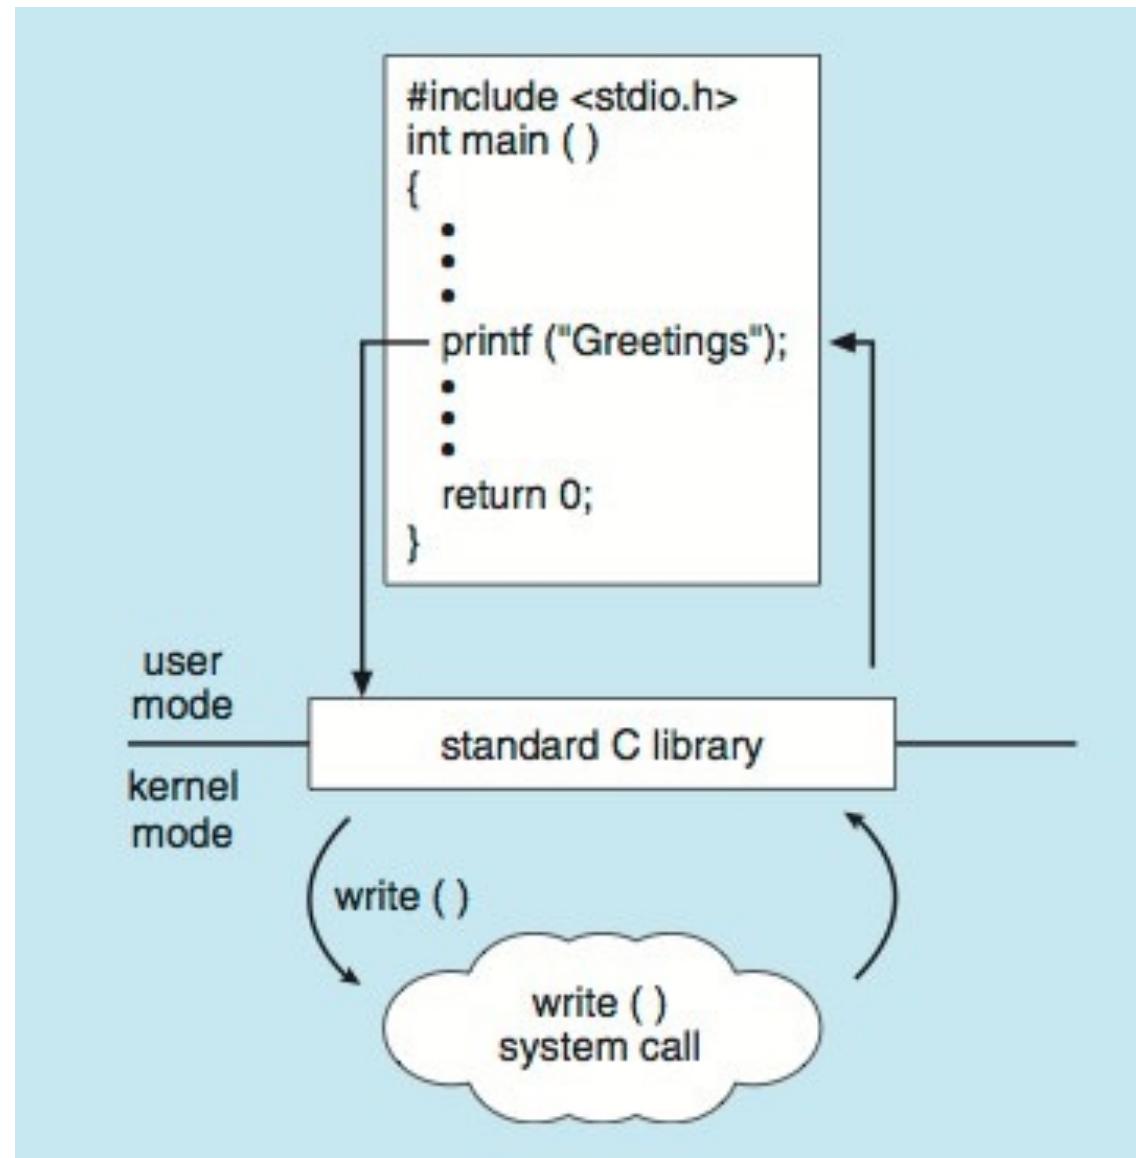
- Programming interface to OS services via system libraries
- Typically written in a high-level language (C, C++)
- Mostly accessed by programs via a high-level *application programming interface* (API) (versus direct system call use)
  - Win32 (Windows), POSIX (Unix, Linux, MacOS), Java (JVM)
- Typically, a number is associated with each system call
  - System-call interface maintains a table indexed by call #
- System call interface invokes the intended system call in OS kernel and returns status of the system call and return values
- Caller just obeys API and understand what OS will do
  - Most details of OS interface hidden by API

***Note: names of systems calls in slides are generic***

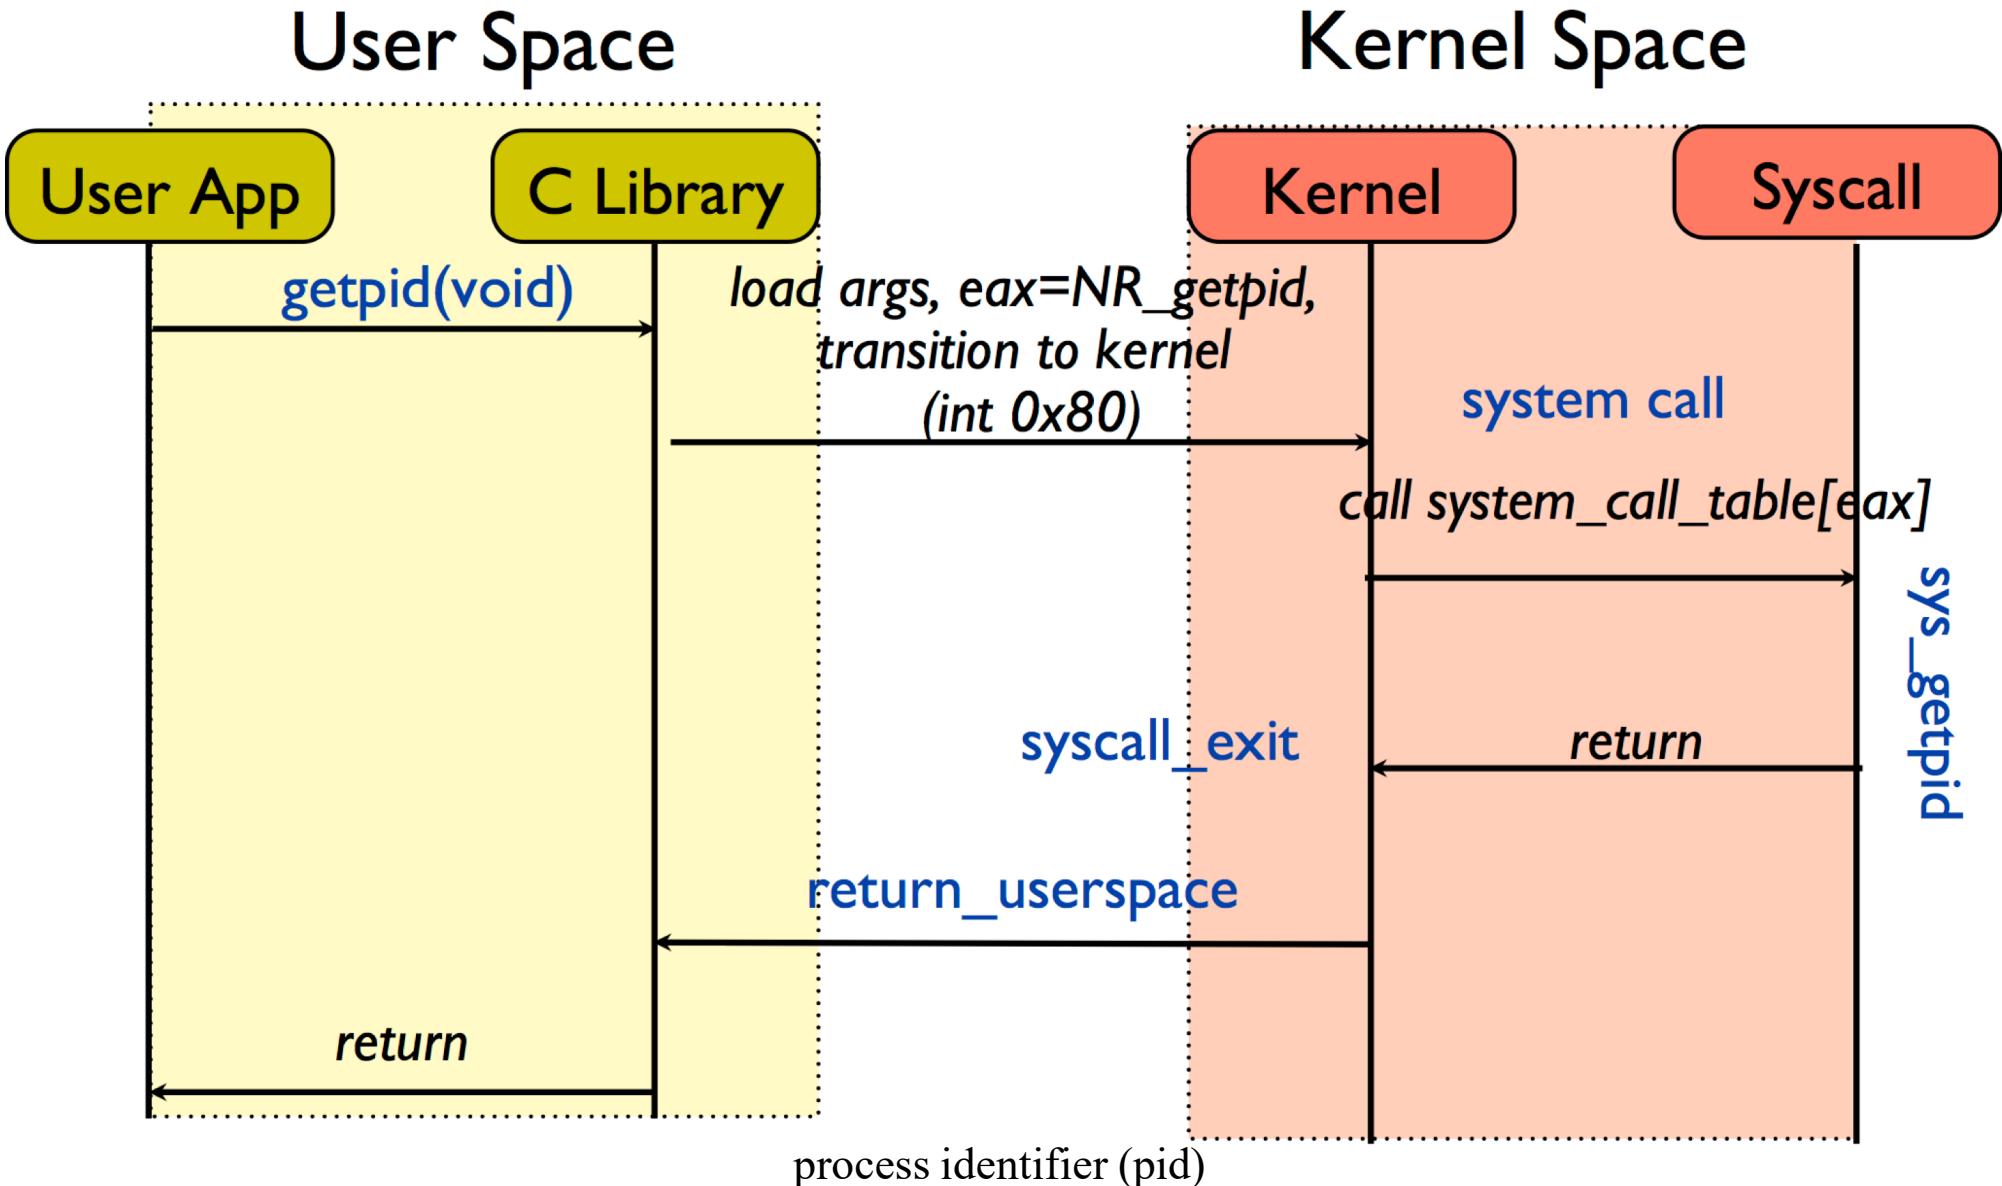
# *System Call – OS Relationship*



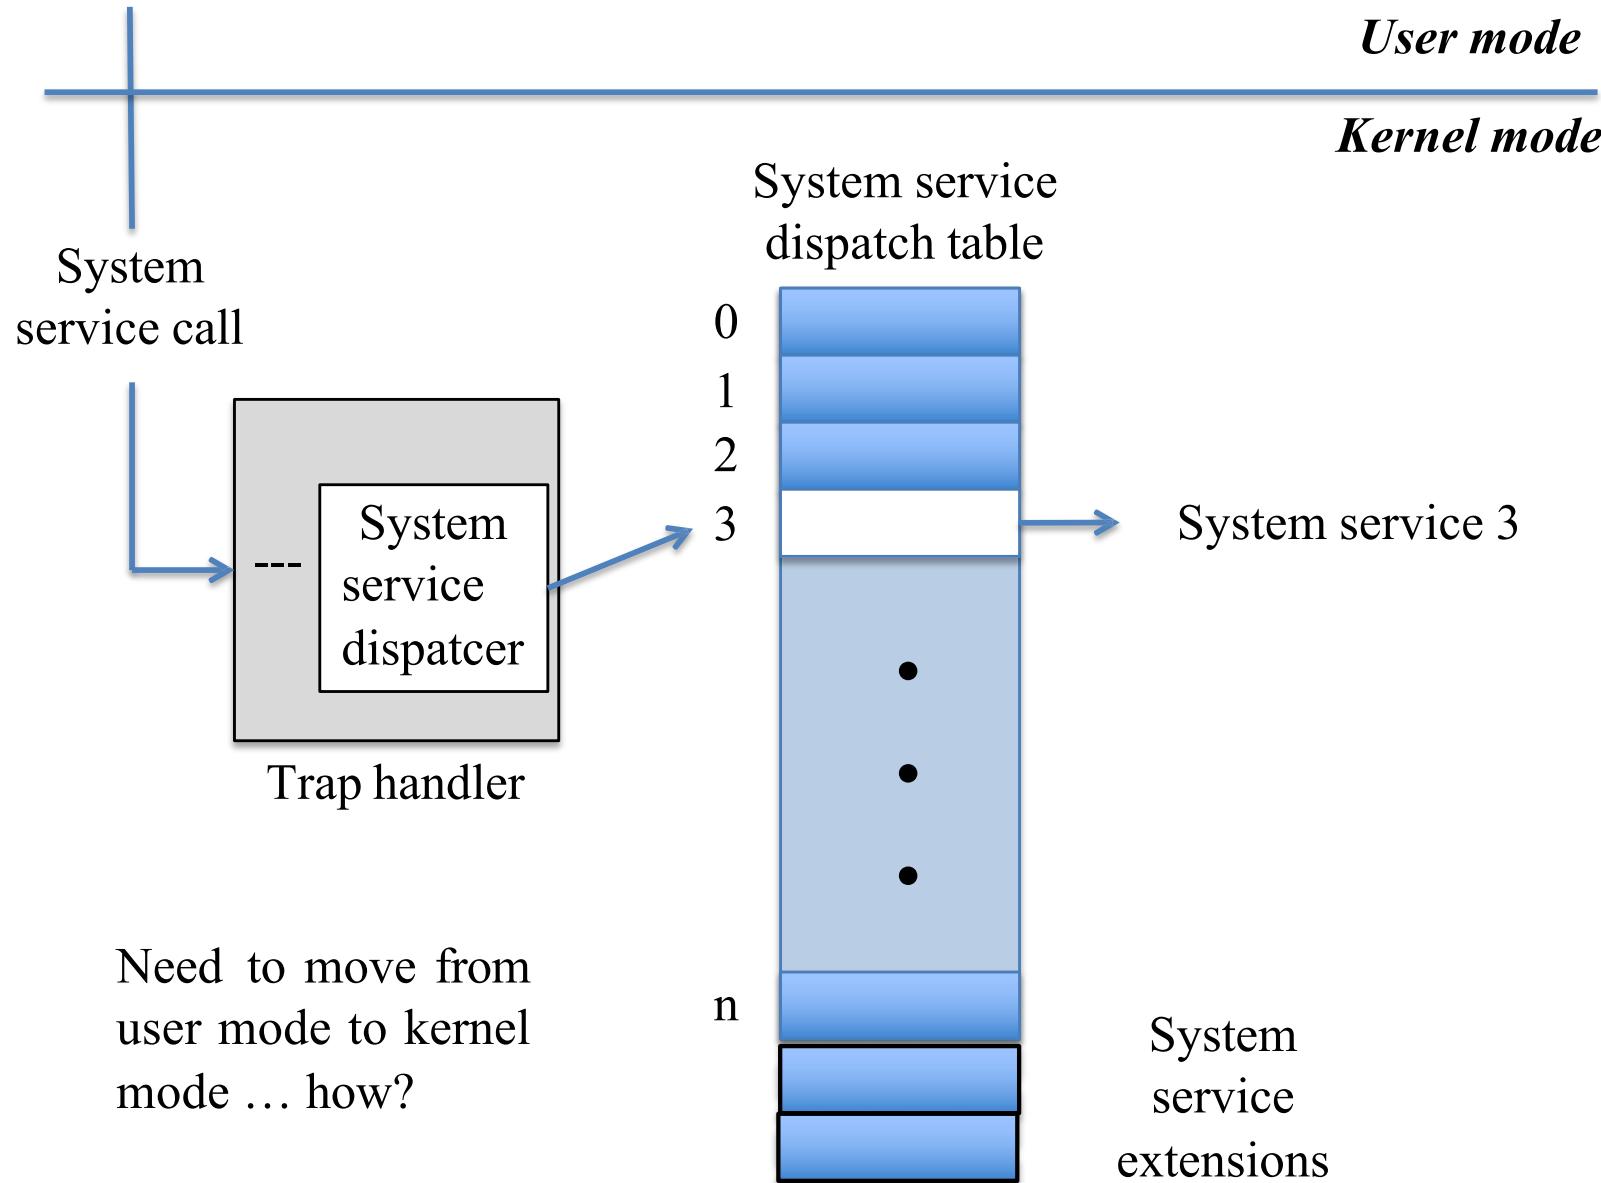
# Standard C Library Example (`printf()`)



# System Call Example (`getpid()`)



# *System Call Handling*

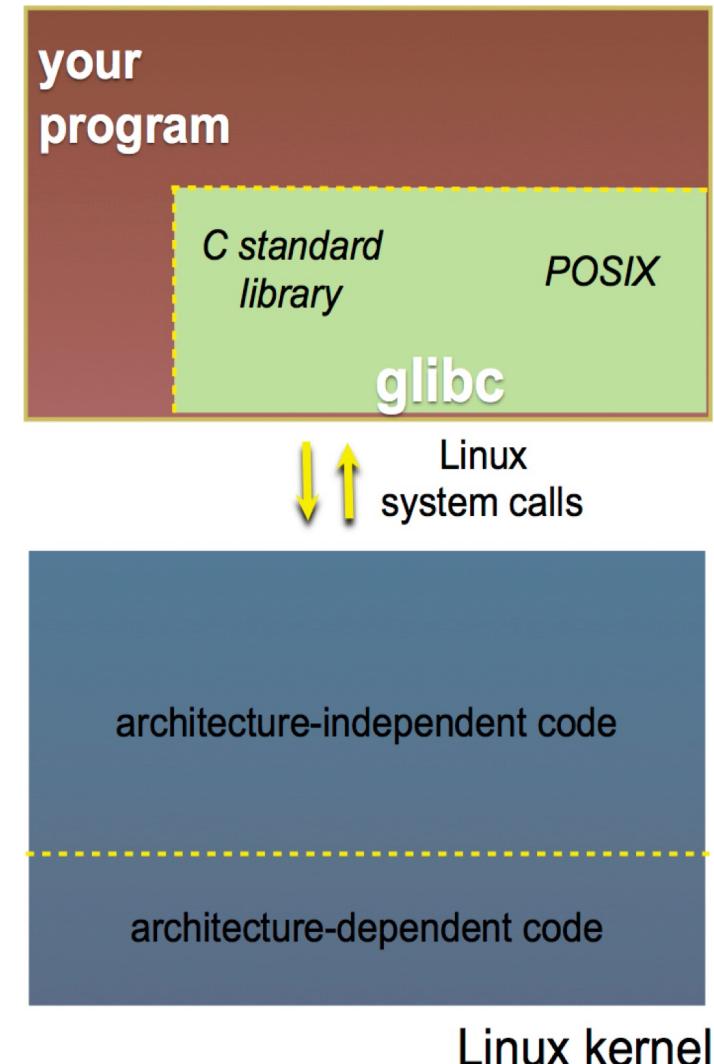


# *System Call Process*

- Procedure call in user process
- Initial work in user mode *libc*
- Trap instruction to invoke kernel *int 0x80*
- Invoke system calls *sys\_read, mmap2*
- Do I/O *read from disk*
- Wait *disk is slow*
- Completion interrupt handling
- Return-from-interrupt instruction
- Final work in library *libc*
- Return to user code

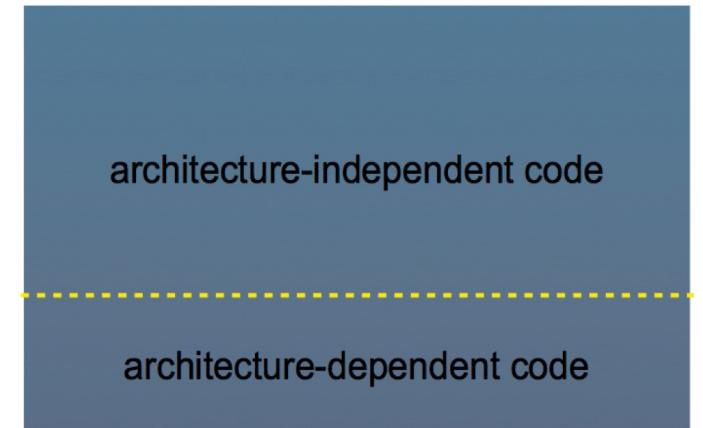
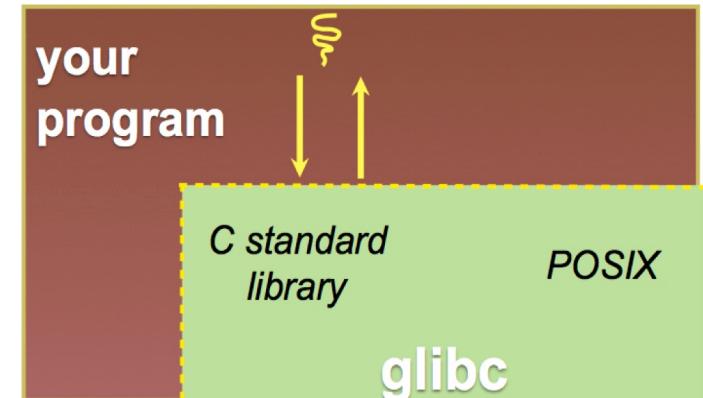
# *Details on x86 / Linux*

- A more accurate picture:
  - Consider a typical Linux process
  - Its course of execution can be in several places
    - ◆ in your program's code
    - ◆ in C standard library
    - ◆ in the kernel model (system calls)



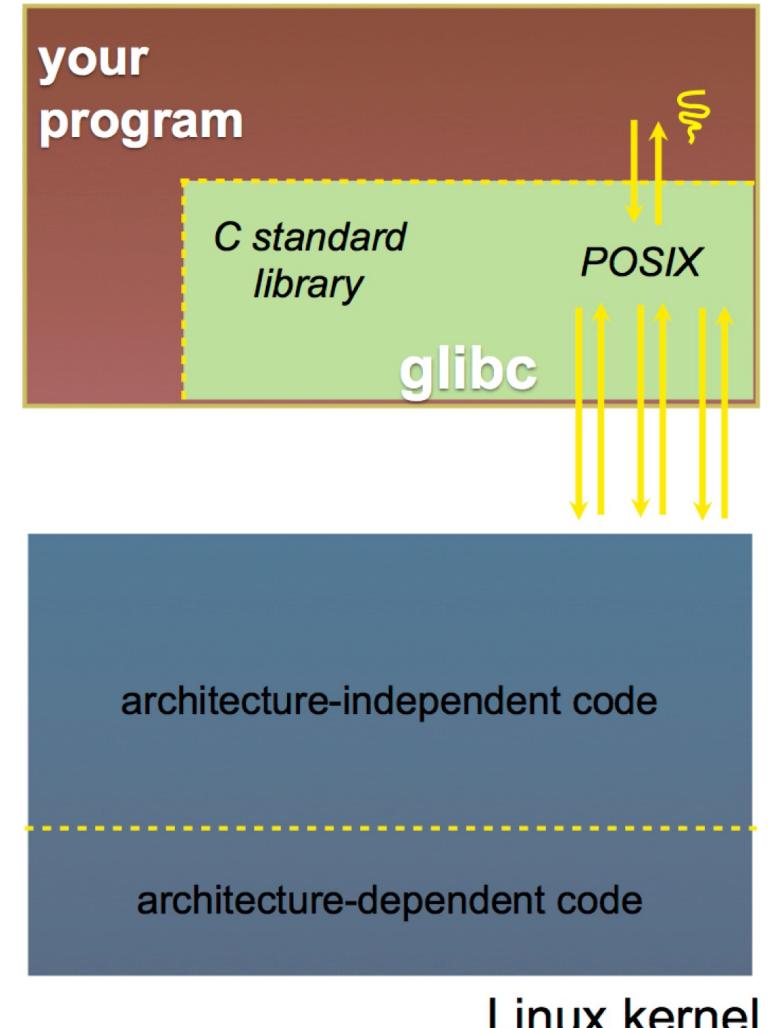
# Details on x86 / Linux

- Some routines your program invokes may be entirely handled by glibc
  - Without involving the kernel◆ for example, *strcmp()* from stdio.h
  - There is some initial overhead when invoking functions in dynamically linked libraries ...
  - ... but after symbols are resolved, invoking glibc routines is nearly as fast as a function call within your program itself



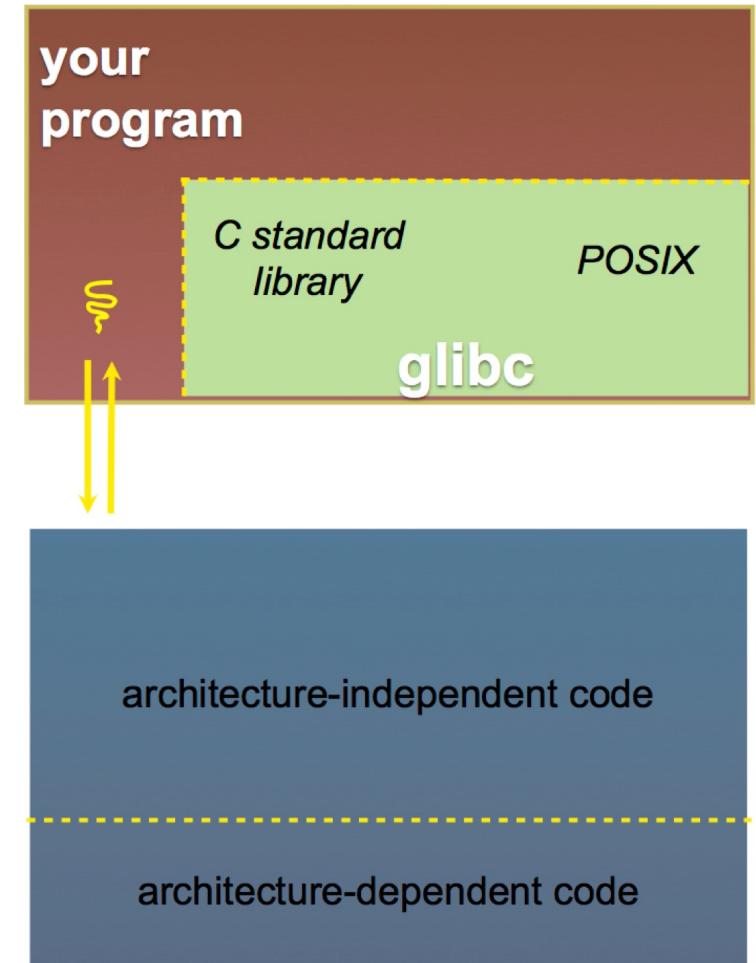
# Details on x86 / Linux

- Some routines may be handled by glibc, but they in turn invoke Linux system calls
  - Example: POSIX wrappers around Linux syscalls
    - ◆ POSIX *readdir()* invokes the underlying Linux *readdir()*
  - Example: C *stdio* functions that read and write from files
    - ◆ *fopen()*, *fclose()*, *fprintf()*, ... invoke underlying Linux *open()*, *read()*, *write()*, *close()*, ...



# Details on x86 / Linux

- Your program can choose to directly invoke Linux system calls as well
  - Nothing forces you to link with glibc and use it
  - But relying on directly invoked Linux system calls may make your program less portable across UNIX varieties



# *Types of System Calls*

- Process control
- File management
- Device management
- Information maintenance
- Communications

# *System Call Wrappers*

- Typically, a program does NOT directly invoke system calls
- Instead, **system call wrappers** are used most of the time
  - Convenience
  - Portability
  - Error handling
  - Abstraction

```
asm volatile (
    "syscall"
    : "=a" (ret)
    : "a"(n), "D"(a1), "S"(a2), "d"(a3)
    : "rcx", "r11", "memory"
);
```

System call instruction

Most of the time, when we say 'system calls', we are actually referring to the system call wrappers.

int fd = open("file.txt", O\_RDONLY);

System call wrapper

# *System Call Wrappers*

<https://man7.org/linux/man-pages/man2/open.2.html>

## open(2) – Linux manual page

[NAME](#) | [LIBRARY](#) | [SYNOPSIS](#) | [DESCRIPTION](#) | [RETURN VALUE](#) | [ERRORS](#) | [VERSIONS](#) |  
[STANDARDS](#) | [HISTORY](#) | [NOTES](#) | [BUGS](#) | [SEE ALSO](#) | [COLOPHON](#)

Search online pages

**open(2)**

System Calls Manual

**open(2)**

**NAME**

[top](#)

open, openat, creat – open and possibly create a file

What does the (2) mean?

How does it differ from fopen?

# *System Call Wrappers vs. C Library Functions*

- Both are in C standard library: glibc
- System call wrapper
  - Thin C function that corresponds almost 1:1 with a kernel system call.
  - Bridge user space and kernel space (for a program to request OS services)
- C Library Functions
  - Higher-level functions
  - Could be purely user-space (string manipulations)
  - Could call system call wrappers to request OS services (open files)

open(2) – Linux manual page

fopen(3) – Linux manual page

(2): System call wrappers  
(3): C library functions

CS 415: Operating Systems, University of Oregon, Fall 2025

UNIVERSITY OF OREGON

# *System Call Wrappers vs. C Library Functions*

Your Program (C code)  
calls read(2), write(2), open(2), etc.



System call wrappers (in glibc)  
• read(2), write(2), open(2)



Kernel (actual syscalls)  
• read(), write(), open()

Directly call system call wrapper

Both are valid ways to make  
system calls

Your Program (C code)  
calls fread(3), printf(3), etc.



libc high-level functions (in glibc)  
• fread(3) calls read(2)  
• printf(3) calls write(2)



System call wrappers (in glibc)  
• read(2), write(2), open(2)



Kernel (actual syscalls)  
• read(), write(), open()

C library functions call system call wrapper

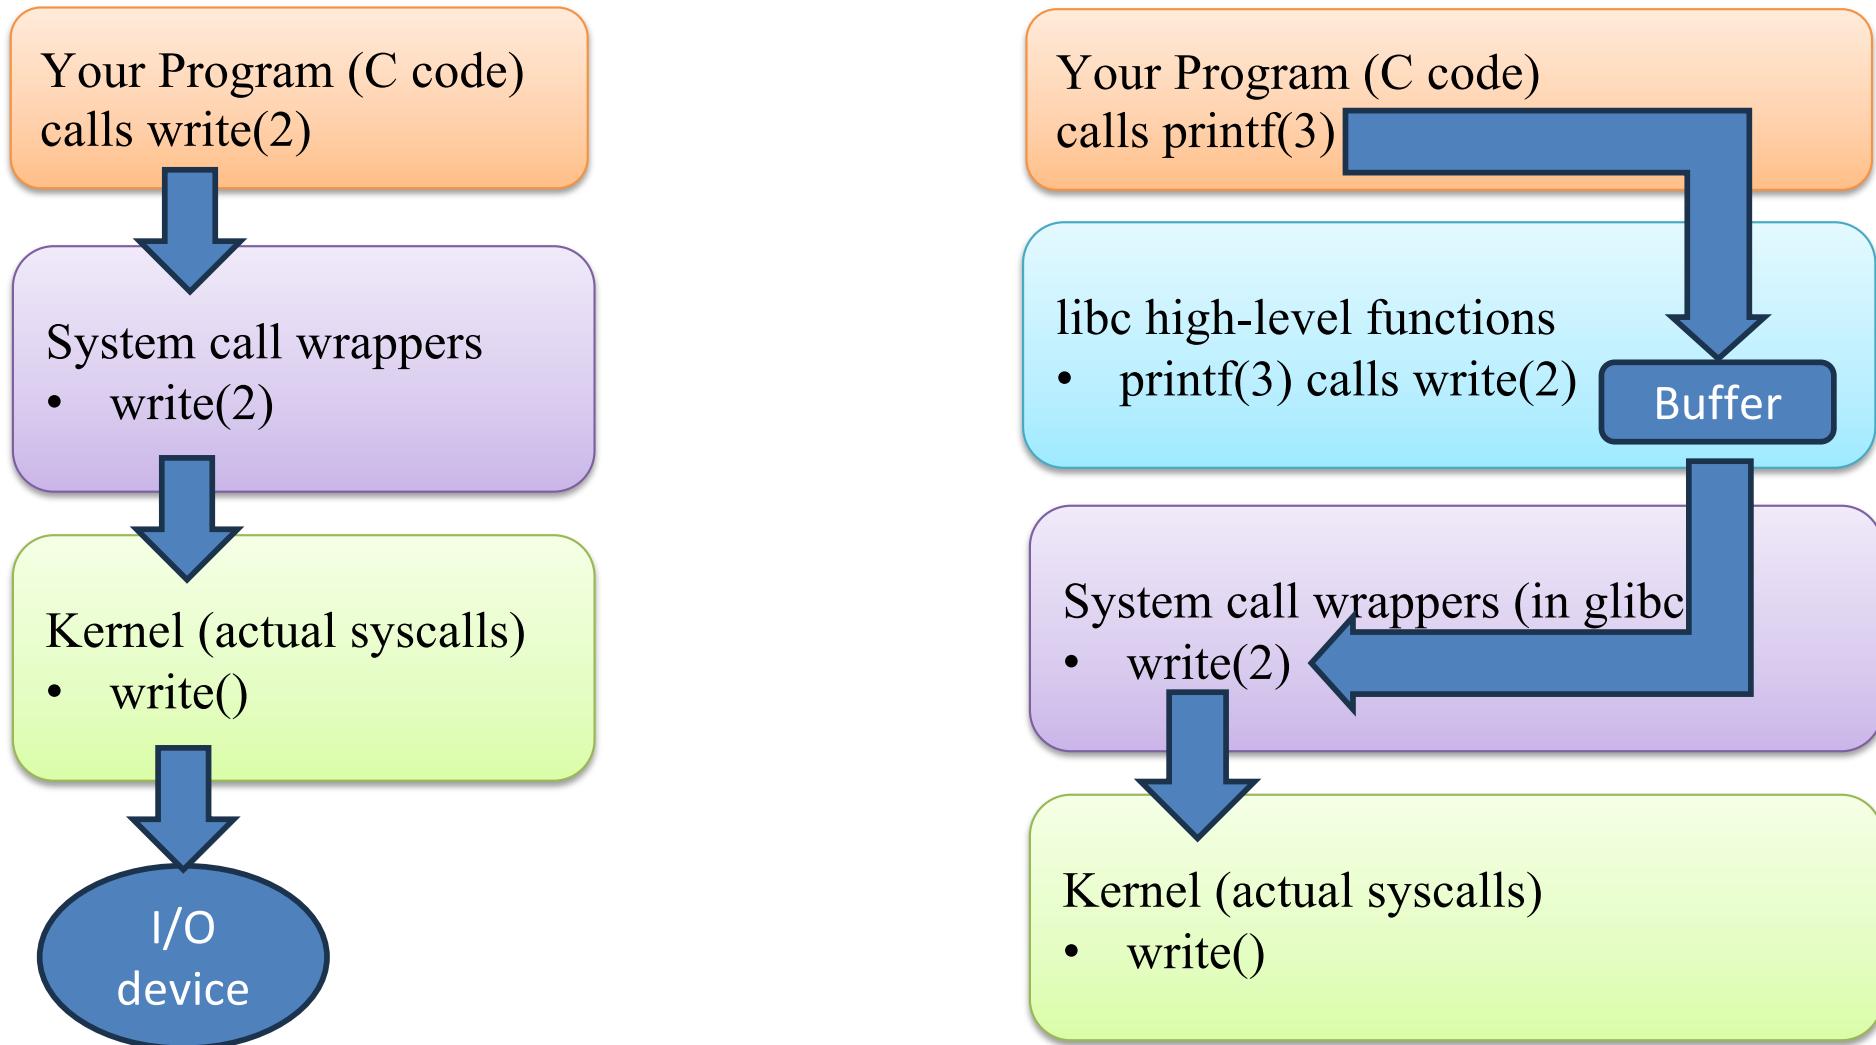
# *System Call Wrappers vs. C Library Functions*

<b>Aspect</b>	<b>System Call Wrapper</b>	<b>C Library Function</b>
<b>Layer</b>	Low-level (full OS service functions)	High-level (less OS service functions)
<b>Location</b>	glibc	glibc
<b>Implementation</b>	Directly maps to a kernel syscall without buffering	May use multiple syscalls internally and may use internal buffers
<b>Performance</b>	Lower latency (immediate response)	More efficient for frequent small I/O operations
<b>Portability</b>	Less portable (depends on OS syscall API)	More portable (follows C/POSIX standards)
<b>Error Handling</b>	Returns -1 and sets errno	Also sets errno, may do extra handling
<b>Example Functions</b>	read(), write(), open(), close(), fork(), execve()	fread(), fwrite(), fopen(), fprintf(), printf(), system(), malloc()

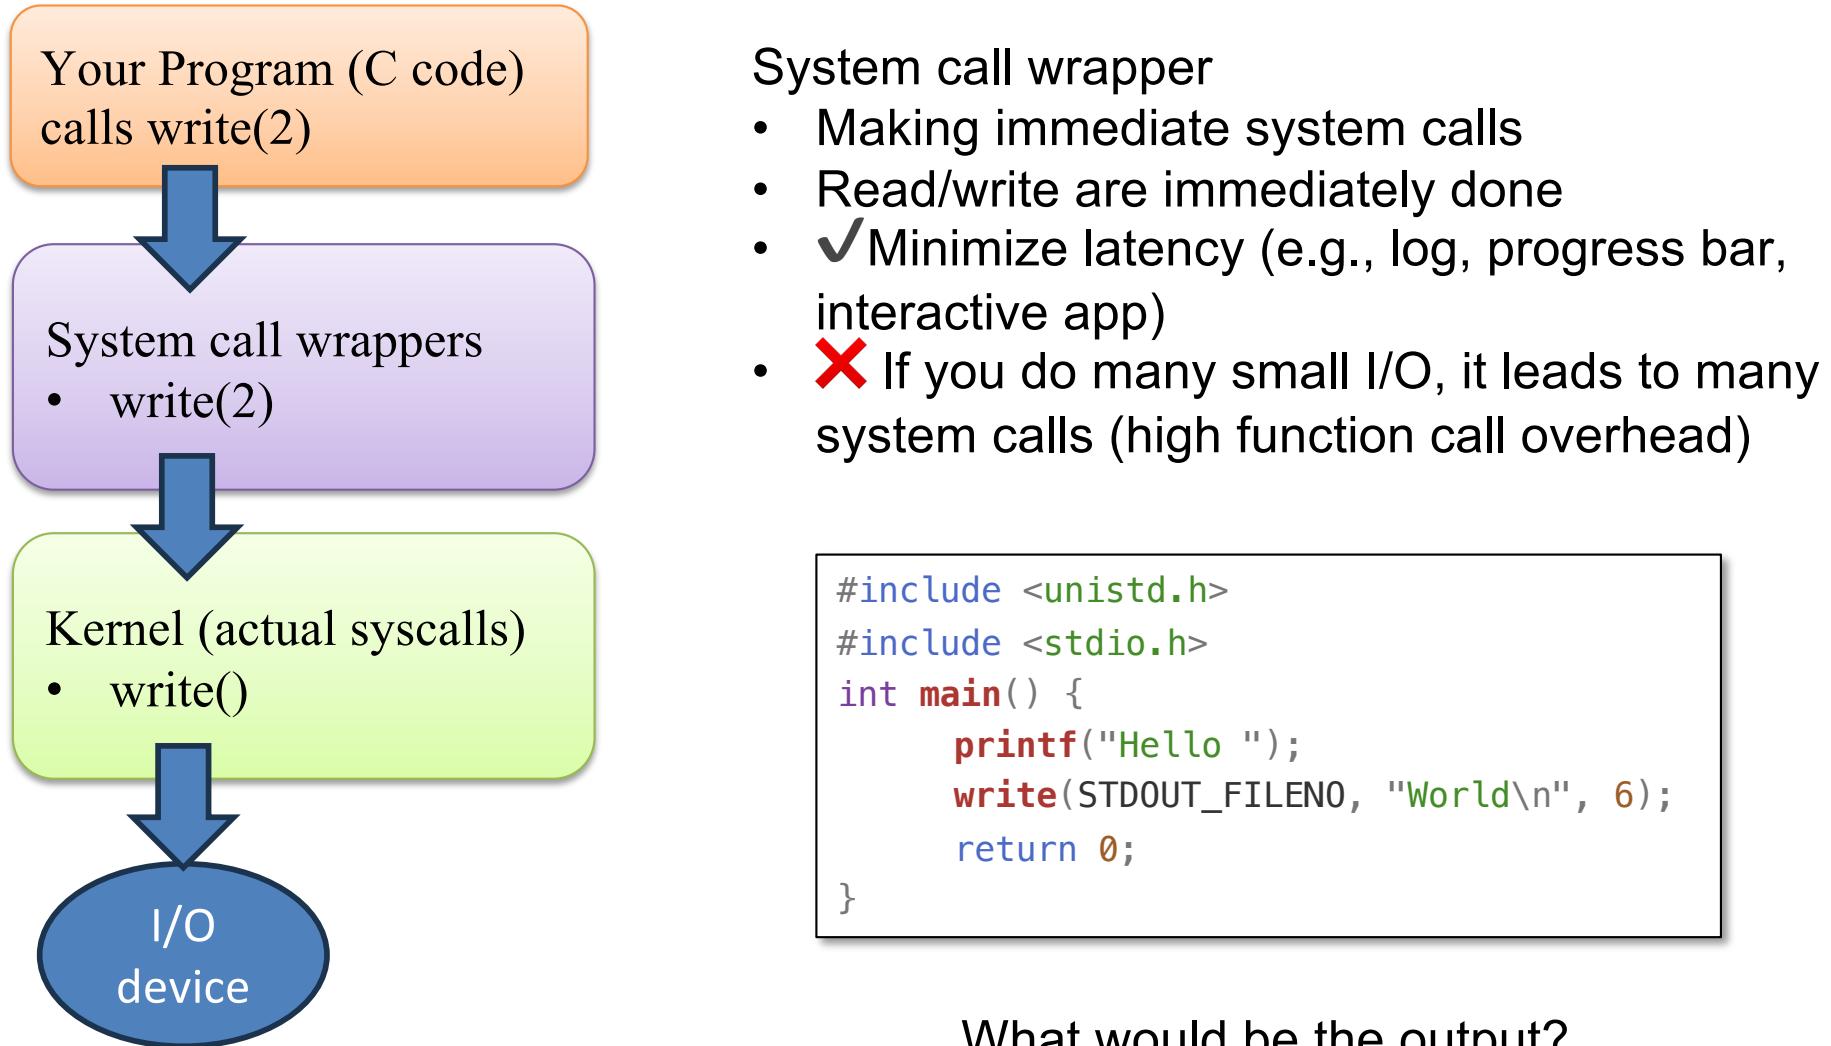
# *System Call Wrappers vs. C Library Functions*

Aspect	System Call Wrapper	C Library Function
Layer	Low-level (full OS service functions)	High-level (less OS service functions)
Location	glibc	glibc
Implementation	Directly maps to a kernel syscall without buffering	May use multiple syscalls internally and may use internal buffers
Performance	Lower latency (immediate response)	More efficient for frequent small I/O operations
Portability	Less portable (depends on OS syscall API)	More portable (follows C/POSIX standards)
Error Handling	Returns -1 and sets errno	Also sets errno, may do extra handling
Example Functions	read(), write(), open(), close(), fork(), execve()	fread(), fwrite(), fopen(), fprintf(), printf(), system(), malloc()

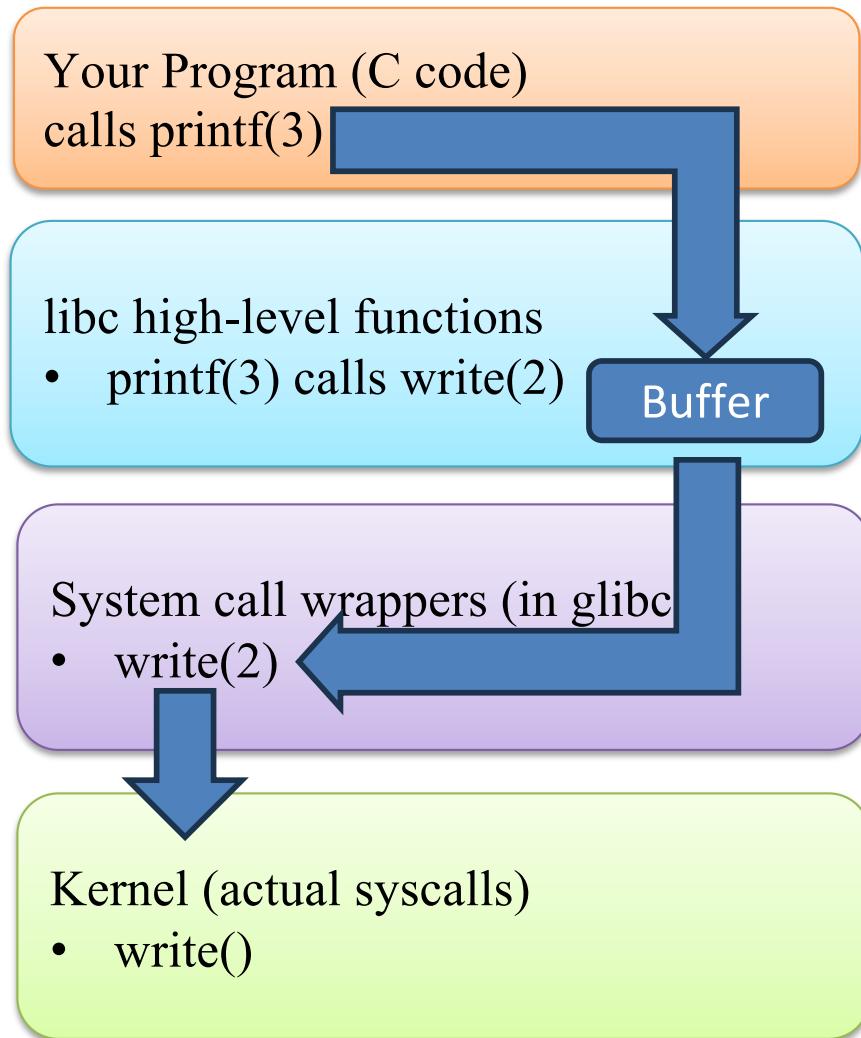
# *System Call Wrappers vs. C Library Functions*



# System Call Wrappers vs. C Library Functions



# System Call Wrappers vs. C Library Functions



## C library function

- Maintain an internal buffer
- Read/write is first done in buffer
- Write out to I/O when buffer is full or we explicitly let it flush
- Read is faster if data is in buffer
- ✓ Minimize the number of system calls
- ✗ May have higher latency (read/write are NOT immediately done)

New line or `fflush(stdout)` can flush stdout buffer

```
#include <unistd.h>
#include <stdio.h>
int main() {
    printf("Hello \n");
    write(STDOUT_FILENO, "World\n", 6);
    return 0;
}
```

What would be the output?

# System Call Wrappers vs. C Library Functions

Your Program (C code)  
calls write(2)

System call wrappers  
• write(2)

Kernel (actual syscalls)  
• write()

I/O device

```
// unbuffered_file_write_timed.c
#include <fcntl.h>
#include <unistd.h>
#include <stdio.h>
#include <time.h>

int main() {
    int fd = open("unbuffered.txt",
O_WRONLY | O_CREAT | O_TRUNC,
0644);
    if (fd < 0) {
        perror("open");
        return 1;
    }

    const char c = 'A';
    for (int i = 0; i < 100000; i++) {
        if (write(fd, &c, 1) != 1) {
            perror("write");
            return 1;
        }
    }

    close(fd);
    return 0;
}
```

Directly call system call  
wrapper

```
// buffered_file_write_timed.c
#include <stdio.h>
#include <time.h>

int main() {
    FILE *fp = fopen("buffered.txt", "w");
    if (!fp) {
        perror("fopen");
        return 1;
    }

    const char c = 'A';
    for (int i = 0; i < 100000; i++) {
        if (fwrite(&c, 1, 1, fp) != 1) {
            perror("fwrite");
            fclose(fp);
            return 1;
        }
    }

    fclose(fp);
    return 0;
}
```

C library functions call  
system call wrapper

How fast can they perform 100,000 small write (1 byte) operations?

# System Call Wrappers vs. C Library Functions

```
// unbuffered_file_write_timed.c
#include <fcntl.h>
#include <unistd.h>
#include <stdio.h>
#include <time.h>

int main() {
    int fd = open("unbuffered.txt",
        O_WRONLY | O_CREAT | O_TRUNC,
        0644);
    if (fd < 0) {
        perror("open");
        return 1;
    }

    const char c = 'A';
    for (int i = 0; i < 100000; i++) {
        if (write(fd, &c, 1) != 1) {
            perror("write");
            return 1;
        }
    }

    close(fd);
    return 0;
}
```

- System call wrappers are more complex  
more functionalities
- C library functions are simpler to use,  
less functionalities

```
// buffered_file_write_timed.c
#include <stdio.h>
#include <time.h>

int main() {
    FILE *fp = fopen("buffered.txt", "w");
    if (!fp) {
        perror("fopen");
        return 1;
    }

    const char c = 'A';
    for (int i = 0; i < 100000; i++) {
        if (fwrite(&c, 1, 1, fp) != 1) {
            perror("fwrite");
            fclose(fp);
            return 1;
        }
    }

    fclose(fp);
}

return 0;
}
```

Directly call system call  
wrapper

- fd and fp are different types
- Both are for referring to a file
- System call wrapper use file descriptor,
  - int type, managed by the kernel
- C library function uses file stream
  - FILE type, managed by C library
- Can convert between each other
  - fdopen ()
  - fileno

C library functions call  
system call wrapper

# *System Call Wrappers vs. C Library Functions*

Aspect	System Call Wrapper	C Library Function
Layer	Low-level (full OS service functions)	High-level (less OS service functions)
Location	glibc	glibc
Implementation	Directly maps to a kernel syscall without buffering	May use multiple syscalls internally and may use internal buffers
Performance	Lower latency (immediate response)	More efficient for frequent small I/O operations
Portability	Less portable (depends on OS syscall API)	More portable (follows C/POSIX standards)
Error Handling	Returns -1 and sets errno	Also sets errno, may do extra handling <code>fread()</code> , <code>fwrite()</code> , <code>fopen()</code> , <code>fprintf()</code> , <code>printf()</code> , <code>system()</code> , <code>malloc()</code>
Example Functions	<code>read()</code> , <code>write()</code> , <code>open()</code> , <code>close()</code> , <code>fork()</code> , <code>execve()</code>	

# *System Call Wrappers vs. C Library Functions*

```
#include <stdio.h>
#include <fcntl.h>
#include <unistd.h>

int main() {
    int fd = open("example.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);
    if (fd == -1) {
        perror("open");
        return 1;
    }

    const char *text = "Hello, world!\n";
    if (write(fd, text, 14) != 14) {
        perror("write");
    }

    close(fd);
    return 0;
}
```

```
#include <stdio.h>

int main() {
    FILE *fp = fopen("example.txt", "w");
    if (!fp) {
        perror("fopen");
        return 1;
    }

    fprintf(fp, "Hello, world!\n");
    fclose(fp);
    return 0;
}
```

✓ Works fine on **Linux, macOS, other POSIX systems.**

✗ **Will fail on Windows:** Windows doesn't use open/write directly  
Windows has CreateFile and other Win32 APIs instead.

✓ Works on **Linux, macOS, Windows.**  
The C standard library implements fopen/fwrite for each OS internally.  
You don't have to worry about the underlying system calls or permissions differences.

# *File Interface*

- Goal:
  - Provide a uniform abstraction for accessing the OS and its resources
- Abstraction:
  - File
- Use file system calls to access OS services
  - Devices, sockets, pipes, and so on
  - Also used in OS in general

# *I/O with System Calls*

- Much I/O is based on a streaming model
  - Sequence of bytes
- *write()* sends a stream of bytes somewhere
- *read()* blocks until a stream of input is ready
- Annoying details:
  - Might fail, can block for a while
  - Working with file descriptors
  - Arguments are pointers to character buffers
  - See the *read()* and *write()* man pages

# *File Descriptors*

- A process might have several different I/O streams in use at any given time
- These are specified by a kernel data structure called a *file descriptor*
  - Each process has its own table of file descriptors
- *open()* associates a file descriptor with a file
- *close()* destroys a file descriptor
- Standard input and standard output are usually associated with a terminal
  - More on that later

# File Descriptors

```
// header files omitted
int main() {
    const char *msg_stdout = "This goes to stdout (terminal)\n";
    const char *msg_stderr = "This is an error message (stderr)\n";
    write(STDOUT_FILENO, msg_stdout, 33);
    write(STDERR_FILENO, msg_stderr, 34);

    int log_fd = open("log.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);
    if (log_fd == -1) {
        perror("open log.txt");
        return 1;
    }
    const char *log_msg = "This is a log message\n";
    write(log_fd, log_msg, 22);

    int input_fd = open("input.txt", O_RDONLY);
    if (input_fd == -1) {
        perror("open input.txt");
        close(log_fd);
        return 1;
    }

    char buffer[100];
    ssize_t n = read(input_fd, buffer, sizeof(buffer) - 1);
    write(1, buffer, n); // write to stdout

    close(log_fd);
    close(input_fd);
}
```

STDOUT\_FILENO=1: standard output  
STDERR\_FILENO=2: standard error

Open a file for writing logs

Open a file for reading input

# Regular File

- File has a *pathname*: `/tmp/foo`
- Can open the file
  - `int fd = open( "/tmp/foo", O_RDWR )`
  - For reading and writing
- Can read from and write to the file
  - `bytes = read(fd, buf, max); /* buf get output */`
  - `bytes = write(fd, buf, len ); /* buf has input */`

*flags for  
read/write  
access*

*pointer to  
buffer*

# Socket File

- File has a pathname: */tmp/bar*
  - Files provide a persistence for a communication channel
  - Usually used for local communication (UNIX domain sockets)
- Open, read, and write via socket operations
  - $sockfd = socket( AF\_INET, SOCK\_STREAM, 0 );$
  - *local.path* is set to */tmp/bar*
  - *bind ( sockfd, &local, len )*
  - Use sock operations to read and write

# Socket File

```
// header files omitted
int main() {
    int sockfd = socket(AF_INET, SOCK_STREAM, 0);
    // create socket
    if (sockfd < 0) {
        perror("socket");
        return 1;
    }

    struct sockaddr_in server_addr;
    memset(&server_addr, 0, sizeof(server_addr));
    server_addr.sin_family = AF_INET;
    server_addr.sin_port = htons(8080); // server port
    if (inet_pton(AF_INET, "127.0.0.1",
&server_addr.sin_addr) <= 0) {
        perror("inet_nton");
        close(sockfd);
        return 1;
    }

    // connect to server
    if (connect(sockfd, (struct sockaddr*)&server_addr,
sizeof(server_addr)) < 0) {
        perror("connect");
        close(sockfd);
        return 1;
    }
}
```

File descriptor for network socket

```
// send data using file descriptor
const char *msg = "Hello server\n";
write(sockfd, msg, strlen(msg));

// receive response
char buffer[1024];
ssize_t n = read(sockfd, buffer, sizeof(buffer)-1);
if (n > 0) {
    buffer[n] = '\0';
    printf("Received: %s", buffer);
}

// close socket
close(sockfd);
return 0;
}
```

Use write/read to send/receive data over networks

# Device File

- Files for interacting with physical devices
  - */dev/null* (no device)
  - */dev/cdrom* (CD-drive)
- Use file system operations, but are handled in device-specific ways
  - *open()*, *read()*, *write()* correspond to device-specific functions (act as function pointers!)
  - Also, use *ioctl* (I/O control) to interact (later)

# Device File

```
// header files omitted
int main() {
    // Open CD-ROM device
    int fd = open("/dev/cdrom", O_RDONLY);
    if (fd < 0) {
        perror("open /dev/cdrom");
        return 1;
    }

    // Read first 2048 bytes (one CD-ROM sector)
    unsigned char buffer[2048];
    ssize_t n = read(fd, buffer, sizeof(buffer));
    if (n < 0) {
        perror("read");
        close(fd);
        return 1;
    }

    printf("Read %zd bytes from CD-ROM\n", n);

    // Close the device
    close(fd);
    return 0;
}
```

Reading from the CD-ROM

# *Other System Calls*

- It's possible to hook the output of one program into the input of another
  - *pipe()*
- It's possible to block until one of several file descriptor streams is ready
  - *select()*
- Special calls for dealing with network
  - AF\_INET sockets, and so on
- Send a message to other (or all) processes
  - *signal()*
- Most of these in section 2 of manual

# *Syscall Functionality*

- System calls are the main interface between processes and the OS
  - Like an extended “instruction set” for user programs that hide many details
  - First Unix system had a couple dozen system calls
  - Current systems have many more
    - ◆ >300 in Linux and >500 in FreeBSD
  - Understanding the system call interface of a given OS lets you write useful programs under it
- Natural questions to ask:
  - Is this the right interface? how to evaluate?
  - How can these system calls be implemented?

# Syscall Functionality

man7.org > Linux > man-pages

Linux/UNIX system programming training

## Linux manual pages: section 2

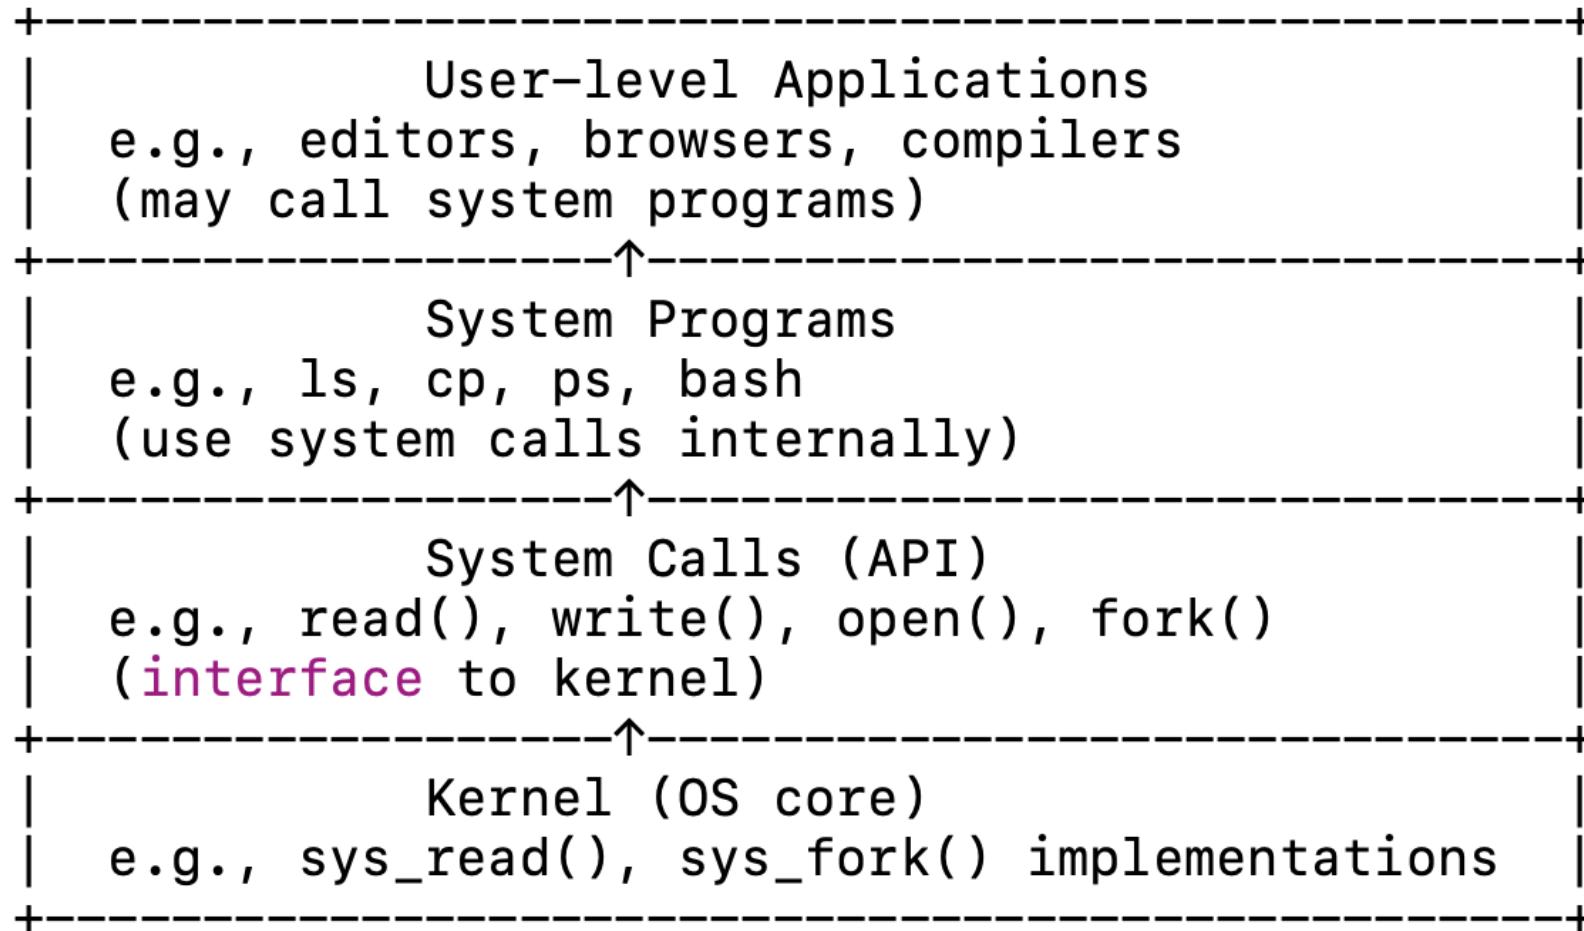
accept(2)	KDSETLED(2const)	PR_SET_UNALIGN(2const)
accept4(2)	kdsettled(2const)	pr_set_unalign(2const)
access(2)	KDSETMODE(2const)	PR_SET_VMA(2const)
acct(2)	kdsetmode(2const)	pr_set_vma(2const)
add_key(2)	KDSIGACCEPT(2const)	PR_SVE_GET_VL(2const)
adjtimex(2)	kdsigaccept(2const)	pr_sve_get_vl(2const)
afs_syscall(2)	KDSKBENT(2const)	PR_SVE_SET_VL(2const)
alarm(2)	kdskbent(2const)	pr_sve_set_vl(2const)
alloc_hugepages(2)	KDSKBLED(2const)	PR_TASK_PERF_EVENTS_DISABLE(2const)
arch_prctl(2)	kdskbled(2const)	pr_task_perf_events_disable(2const)
arm_fadvise(2)	KDSKBMETA(2const)	PR_TASK_PERF_EVENTS_ENABLE(2const)
arm_fadvise64_64(2)	kdskbmeta(2const)	pr_task_perf_events_enable(2const)
arm_sync_file_range(2)	KDSKBMODE(2const)	pselect(2)
bdfflush(2)	kdskbmode(2const)	pselect6(2)
bind(2)	KDSKBSENT(2const)	ptrace(2)
bpf(2)	kdskbsent(2const)	putmsg(2)
break(2)	kexec_file_load(2)	putpmsg(2)
brk(2)	kexec_load(2)	pwrite(2)
cacheflush(2)	keyctl(2)	pwrite64(2)
cachestat(2)	KEYCTL_ASSUME_AUTHORITY(2const)	pwritev(2)
capget(2)	keyctl_assume_authority(2const)	pwritev2(2)
capset(2)	KEYCTL_CHOWN(2const)	query_module(2)
chdir(2)	keyctl_chown(2const)	quotactl(2)
chmod(2)	KEYCTL_CLEAR(2const)	read(2)
chown(2)	keyctl_clear(2const)	readahead(2)
chown32(2)	KEYCTL_DESCRIBE(2const)	readdir(2)
chroot(2)	keyctl_describe(2const)	readlink(2)
clock_adjtime(2)	KEYCTL_DH_COMPUTE(2const)	readlinkat(2)
clock_getres(2)	keyctl_dh_compute(2const)	readv(2)
clock_gettime(2)	KEYCTL_GET_KEYRING_ID(2const)	reboot(2)
clock_nanosleep(2)	keyctl_get_keyring_id(2const)	recv(2)
clock_settime(2)	KEYCTL_GET_PERSISTENT(2const)	recvfrom(2)
clone(2)	keyctl_get_persistent(2const)	recvmsg(2)
clone2(2)	KEYCTL_GET_SECURITY(2const)	remap_file_pages(2)
__clone2(2)	keyctl_get_security(2const)	removexattr(2)
clone3(2)	KEYCTL_INSTANTIATE(2const)	rename(2)
close(2)	keyctl_instantiate(2const)	renameat(2)
close_range(2)	KEYCTL_INSTANTIATE_IOV(2const)	renameat2(2)
connect(2)	keyctl_instantiate iov(2const)	request_key(2)
copy_file_range(2)	KEYCTL_INVALIDATE(2const)	restart_svcall(2)
creat(2)	kevctl_invalidate(2const)	

... and a lot more

# *System Programs*

- System programs provide a convenient environment for program development and execution
- They can be divided into categories:
  - File manipulation
  - Status information
  - File modification
  - Programming language support (program development)
  - Program loading and execution
  - Application programs
- Most user's view of the operating system is defined by system programs, not the actual system calls

# *System Programs vs System Calls*



\* Generated by ChatGPT



UNIVERSITY  
OF OREGON

# Linux man pages online

The links from this page display HTML renderings of the man pages from the Linux [man-pages](#) project as well as a [curated](#) collection of pages from various other free software projects.

- List of all man pages: [by section](#) | [alphabetically](#) | [by project](#)
- Individual sections (page names only): [1](#), [2](#), [3](#), [4](#), [5](#), [6](#), [7](#), [8](#)
- Intro pages: [intro\(1\)](#), [intro\(2\)](#), [intro\(3\)](#), [intro\(4\)](#), [intro\(5\)](#), [intro\(6\)](#), [intro\(7\)](#), [intro\(8\)](#)

Section	Contents	Example
1	System programs (User commands)	ls(1), grep(1)
2	System calls (functions provided by the kernel)	read(2), open(2)
3	Library functions (C standard library, other libraries)	printf(3), malloc(3)
4	Special files / devices	/dev/null(4), /dev/sda(4)
5	File formats and conventions	/etc/passwd(5), crontab(5)
6	Games and screensavers	nethack(6)
7	Miscellaneous / conventions / standards	signal(7), regex(7)
8	System administration commands (root/admin)	mount(8), iptables(8)



# *OS Design and Implementation*

- Design and implementation of OS is not “solvable” in full, but some approaches have proven successful
- Internal structure of different operating systems can vary widely
- Start the design by defining goals and specifications
- Affected by choice of hardware, type of system
- There are *user* goals and *system* goals
  - User goals – operating system should be convenient to use, easy to learn, reliable, safe, and fast
  - System goals – operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient

# *OS Policy versus Mechanism*

- Important principle to separate  
*Policy*: What will (should) be done?  
*Mechanism*: How to do it?
    - Mechanisms determine how to do something
    - Policies decide what will be done
  - Separation of policy from mechanism is a **VERY** important principle
    - Allows maximum flexibility if policy decisions are to be changed later
    - Universal principle
  - Specifying and designing an OS is a highly creative task of software engineering

# *Summary*

- Operating systems must balance many needs
  - Impression that each process has individual use of system
  - Comprehensive management of system resources
- Operating system structures try to make use of system resources straightforward
  - Libraries
  - System services
  - System calls and other interfaces

# *Next Class*

- Processes