



# CS 415

# Operating Systems

## Threads

Prof. Jieyang Chen

Department of Computer and Information Science

Fall 2025



UNIVERSITY OF OREGON

# *Logistics*

---

- Keep working hard on Project 1
  - Due 11:59pm on Friday
  - Please work individually
- Labs this week focus on answering questions
- Take advantage of TA office hours
  - Keep checking Canvas announcements
- Read Chapter 4 on threads
- Encourage you to come to office hours
  - Few people are coming now

# *Outline*

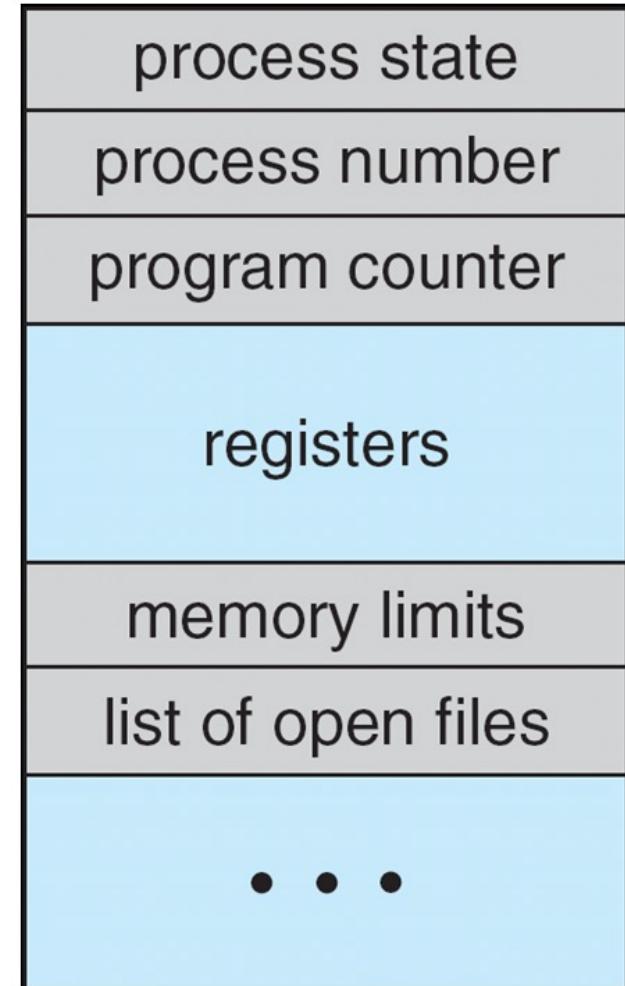
---

- Processes revisited
- Threads

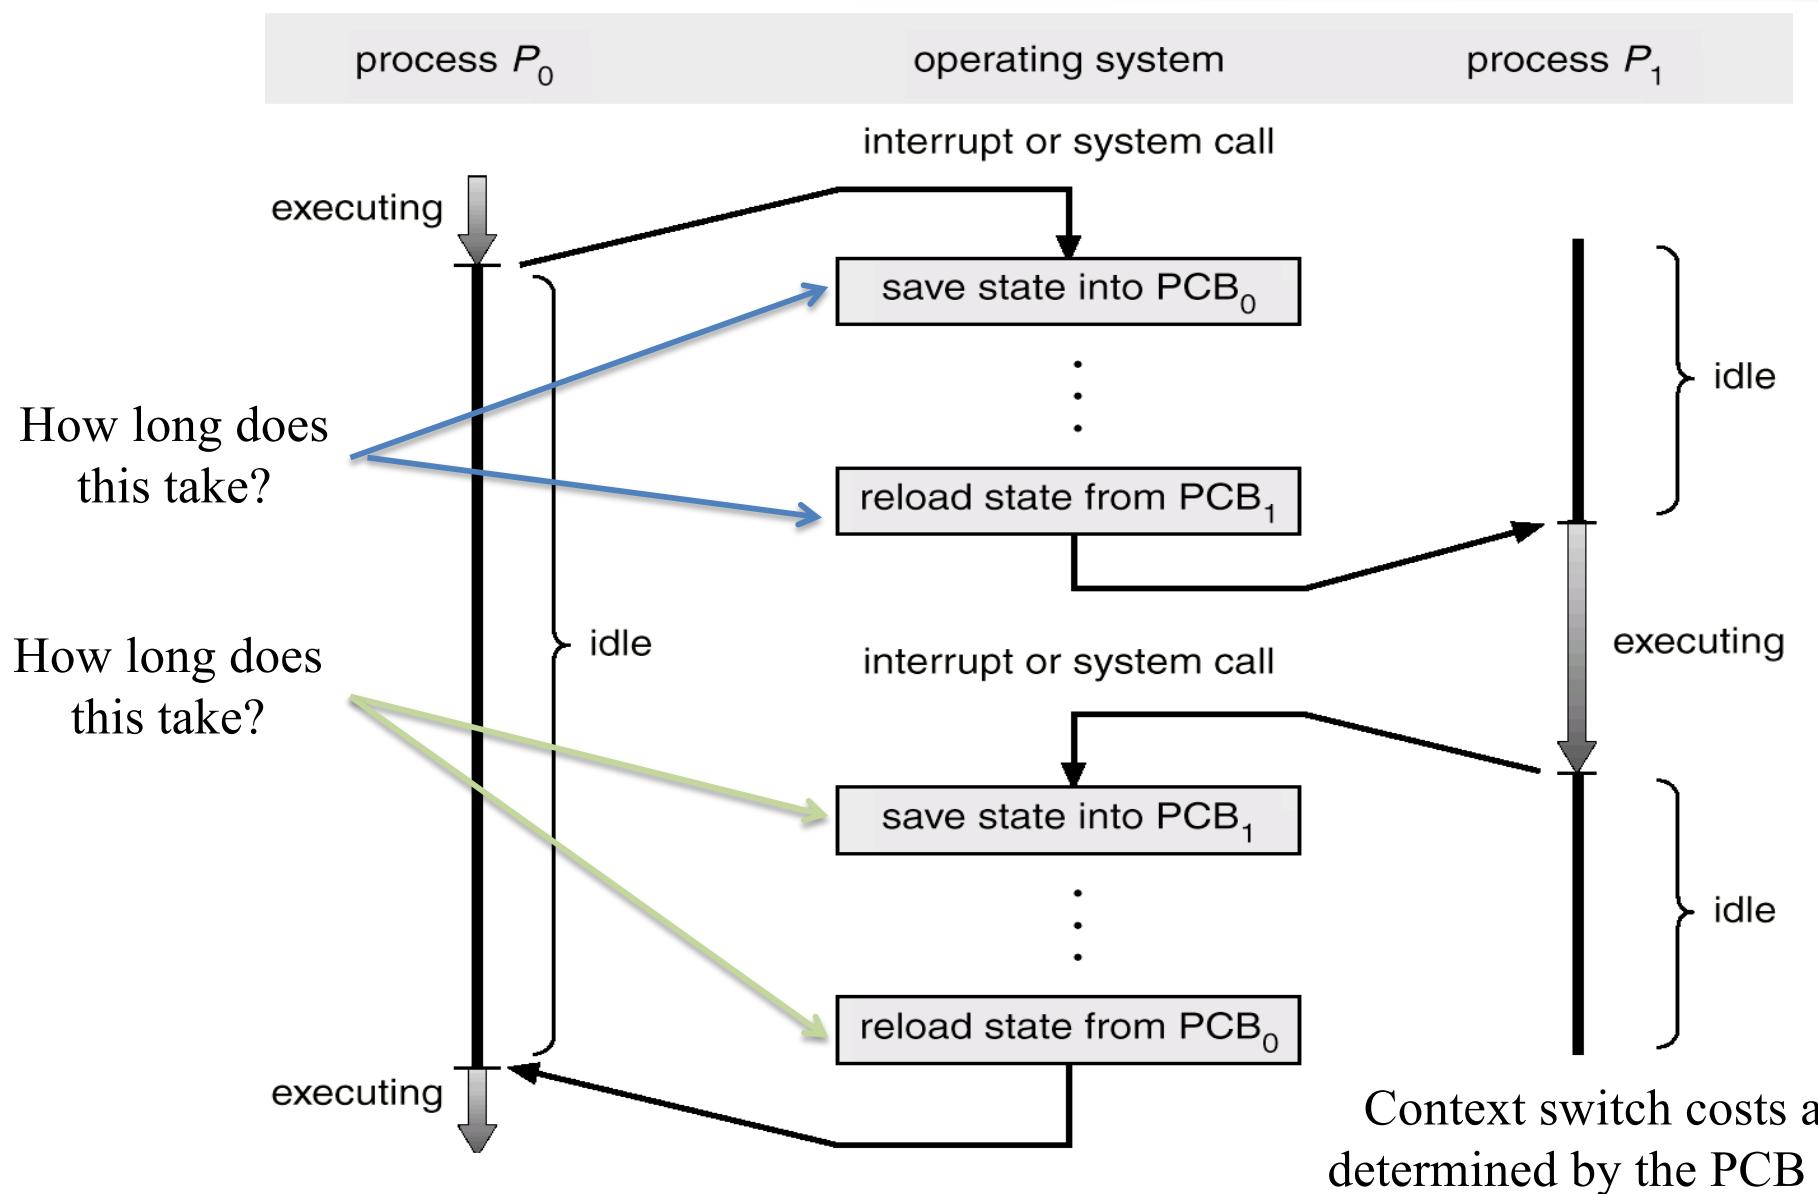
# *Process Control Block*

- Information associated with each process (aka *task control block*)
- Process state
  - Running, waiting, ...
- Program counter
  - Location of instruction to next execute
- CPU registers
  - Contents of all process-centric registers
- CPU scheduling information
  - Priorities, scheduling queue pointers, ...
- Memory-management information
  - Memory allocated to the process
  - Memory limits and other associated information
- Accounting information
  - CPU used, elapsed time, time limits, ...
- I/O status information
  - I/O devices allocated to process, list of open files

**PCB**



# Process Context Switch

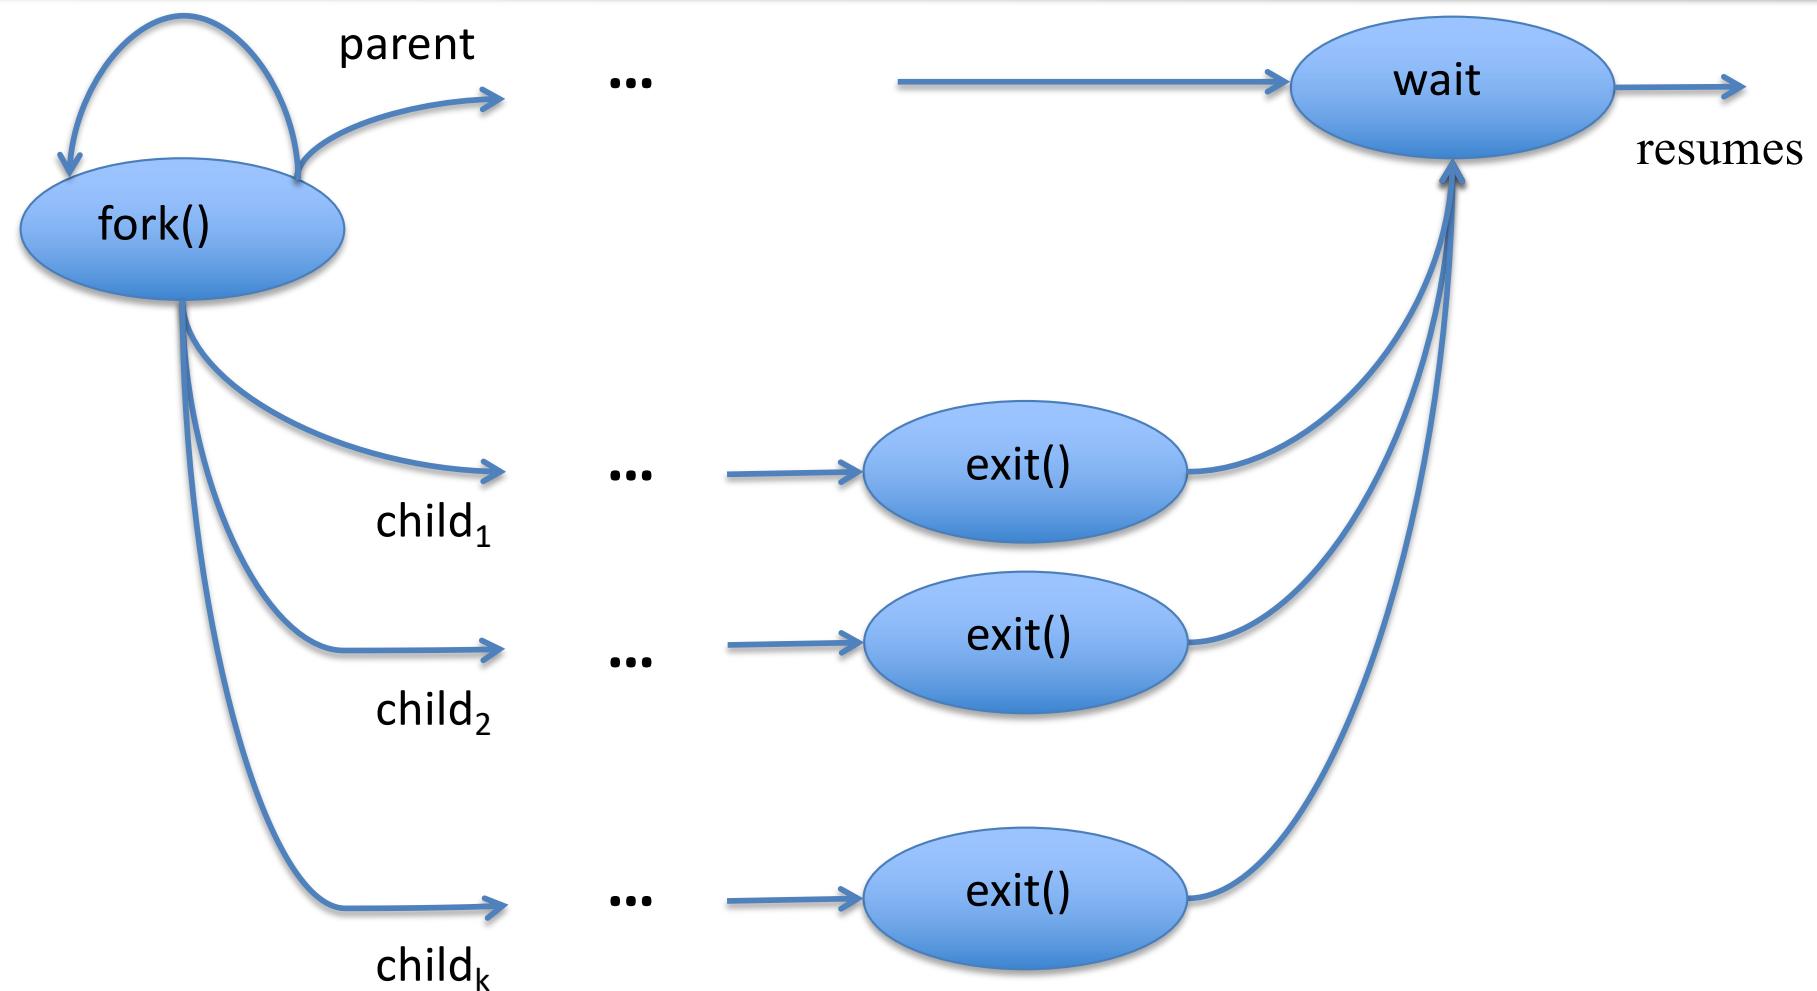


# *Process Model*

---

- Much of the OS's job is keeping processes from interfering with each other ... Why?
- Each process has its OWN resources to use
  - Program code to execute, address space, files, ...
- Processes are good for *isolation (protection)*
  - Prevent one process from affecting another process
- Processes are *heavyweight* ... Why?
  - Pay a price for isolation
  - A full “process swap” is required for multiprocessing
  - There is lots of process state to save and restore
  - OS must context switch between them
    - ◆ intervene to save/restore all process state
- Is there an alternative?

# *Program with Multiple Processes*



# *Program Creation System Calls*

- *fork()*
  - Copy address space of parent and data (copy on write) to child process
- *vfork()*
  - Share the parent's address space
  - Shares address space AND physical memory between parent and child
  - Parent blocks and waits for child to exit or call *exec()*
- *exec()*
  - Load new program and replace address space (for new child code)
  - Some resources may be transferred (open file descriptors)
  - Specified by arguments
- *clone()*
  - Like *fork()* BUT child shares some process context
  - More explicit control over what is shared
  - Calls a function pointed to by argument
  - Process address space and physical memory can be shared!

Child process is  
independently executing

Child process is like  
in *fork()*, but it is  
much lighterweight!

# *Process Reimagined*

---

- A process is “a program in execution”
  - Memory address space containing code and data
  - State information (PC, register, SP) => PCB details
  - Allocated physical memory
  - Other resources (e.g., open file descriptors)
- Consider a process in 2 respects (categories)
  - 1) A collection of resources required for execution
    - ◆ code, address space, memory, open files, ...
  - 2) A “*thread of execution*”
    - ◆ current state of execution (CPU state)
    - ◆ where the execution is now (PC)
- Suppose we think about these separately!
  - Resources
  - Execution state



# *Terminology*

---

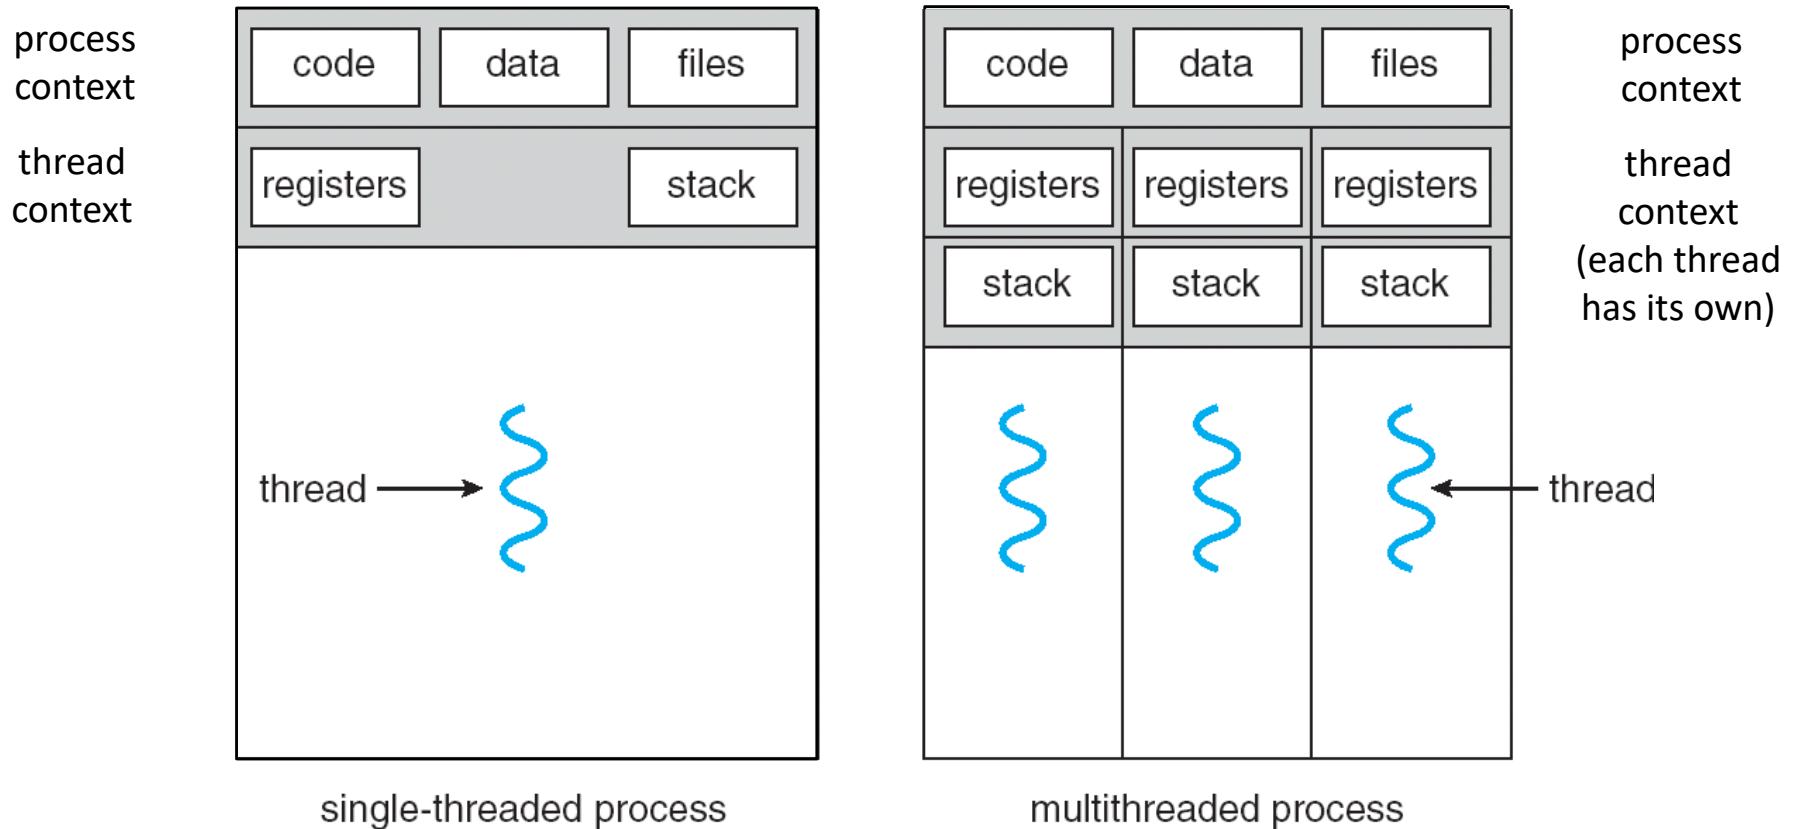
- ***Multiprogramming***
  - Running *multiple programs* on a computer at the same time
    - ◆ Concurrent (logically) execution
  - Each program is a process (or set of processes)
  - Processes of different programs are independent
- ***Multiprocessing***
  - Running *multiple processes* on a computer at the same time
  - Physical concurrently
  - OS schedules processes on processor(s)
- ***Multithreading***
  - Multiple *threads of execution (execution contexts)* in a single address space (of a process)
  - OS executes threads on processor(s)
  - Can be logical or physical concurrent

# *What Makes up a (Multithreaded) Process?*

- Thread of execution (sequence of instructions) on a CPU
  - Program counter *per thread*
  - Registers *per thread*
- Memory
  - Address space (code and data)
    - ◆ address space and physical memory assigned is shared *process*
  - Stack
    - ◆ each thread has its own stack pointer
    - ◆ separately located in the address space*per thread*
  - Heap
    - + private dynamic memory
    - ◆ separately located in the address space*process*  
*per thread*
- I/O
  - Share files, sockets, ... *process*
- PCB
  - A thread lives in a process and uses the PCB *process*

# *Single-Threaded vs. Multithreaded*

- A normal process is a special case of a multithreaded process
  - A process that contains **just one thread!**
- A multithreaded process has multiple threads
  - Where do they live? How do they get created?



# *Why Multithreaded Applications?*

---

- Multiple threads share a common address space
  - More correctly, they share process resources, including memory!
- Why would you want to do that?
- Concurrency!
  - How do you get concurrency?
  - One way is to create multiple processes
  - Use IPC to support process-level concurrency
  - Issues?
- Some applications want to share data structures among concurrently executing parts of the computation
  - Is this possible with processes? => shared segments
  - Is it difficult with processes? => well, it is not that easy
  - Again, use IPC for process-level data sharing
- What is the problem? What is the solution?

# *Advantages of Threads*

---

- Suppose we do not care if the OS enforces resource separation
  - This is a trust issue, in part (back off on isolation / protection)
  - Will introduce issues with respect to concurrency ... What kind?
- Improve responsiveness
  - Possible to have a thread of execution that just responds to events
- Resource sharing is facilitated
  - All threads in a process have equal access to resources
- Economy of resources
  - Thread-level resources are “cheaper” than process resources!
  - Threads are “lighter weight”
- Utilization of multiprocessors
  - Run multiple threads on multiple processes without the overhead of running multiple processes

# *Working with Threads*

---

- In a C program
  - *main()* procedure defines the first thread
  - C programs always start at *main()*
- Now you want to create a second thread
  - Allocate resources to maintain a second execution context in **the same address space**
    - ◆ think about what state will be needed for a thread
  - Want something similar to *fork()* but simpler
    - ◆ supply **a function** when starting the new thread
  - Remember this creates another thread of execution

# *Can a thread call exec() to run a new program?*

---

- Yes, but the results are often NOT what you want
- exec() replaces the entire current process image with a new program.
  - All previous code, data, heap, stack, and threads are destroyed.
  - All other threads are terminated immediately.
  - Only the thread that called exec() survives long enough to run the new program's entry point (main()).
  - The process doesn't become “multi-threaded” anymore — it becomes the new program, running in a single thread.

# *Threads vs. Processes*

---

- Easier to create than a new process
- Less time to terminate a thread than a process
- Less time to switch between two threads
  - Within the same process
- Less communication overheads
  - Communicating between the threads of one process is simple because the threads share everything
  - Address space is shared ...
  - ... thus memory is shared

# *Which is Cheaper?*

- Create new process or create new thread (in existing process)?
- Context switch between processes or threads?
- Interprocess or interthread communication?
- Sharing memory between processes or threads?
- Terminating a process or terminating a thread (not the last one)?

Process creation method	Time (sec), elapsed (real)
<i>fork()</i>	22.27 (7.99)
<i>vfork() (faster fork)</i>	3.52 (2.49)
<i>clone()</i>	2.97 (2.14)

Time to create 100,000 processes (Linux 2.6 kernel, x86-32 system)

*clone()* creates a lightweight Linux process (thread)

# *Implications?*

---

- Consider a web server on a Linux platform
- Measure 0.22 ms per *fork()*
  - Maximum of  $(1000 / 0.22) = 4545.5$  connections/sec
  - 0.45 billion connections per day per machine
    - ◆ fine for most servers
    - ◆ too slow for a few super-high-traffic front-line web services
- Facebook serves  $O(750 \text{ billion})$  page views per day
  - Guess ~1-20 HTTP connections per page
  - Would need 3,000 – 60,000 machines just to handle *fork()*, without doing any work for each connection!
- What is the problem here?

# *Thread Attributes*

- Global to process:
  - memory
  - PID, PPID, GID, SID
  - controlling term
  - process credentials
  - record locks
  - FS information
  - timers
  - resource limits
  - and more...
- Local to specific thread:
  - thread ID
  - CPU state and stack
  - signal mask
  - thread-specific data
  - alternate signal stack
  - error return value
  - scheduling policy/priority
  - Linux-specific  
(e.g., CPU affinity)

thread state is small  
versus  
process state  
(stored in PCB)

# *Thread Programming and Execution*

---

- Every process is created with 1 “thread of execution”
  - Starts executing program at main()
- Additional threads are created using a threading programming library
  - A new “thread of execution” is created
  - Thread runs in the context of the process
  - Multiple threads can be created
- A process is “executing” when 1 or more of its threads are executing
- Suppose there are  $N$  threads and only 1 CPU
  - 1 thread can be running
  - what’s going on with the other  $N-1$  threads?

# *Kernel Threads*

---

- There is thread management support in the OS
  - Kernel operations for executing process threads
  - Thread “objects” in the kernel to hold thread context
  - Called *kernel threads*
- A process thread must have a kernel thread to run
  - Kernel thread acts like an “**execution container**”
- Nearly all OSes support a notion of threads
  - Linux -- thread and process abstractions are mixed
  - Mac OS X, Windows XP, ...

# *Something Subtle Going on Here (1)*

---

- A thread is just a (logical) unit of concurrency
  - It has a stack, a PC, a CPU state, and ... that's it!
  - It is a “thread of execution” in the process code
- A thread can “execute” when it is assigned to a “resource” which is then scheduled on a CPU
  - That resource is a **kernel thread**
- Where is the thread when it is not executing?
  - Two parts
    - ◆ (user) thread control block (TCB) + kernel thread control block (KTCB)
  - Condition 1: no kernel thread assigned
    - ◆ Only TCB is saved in the memory of the process!
    - ◆ KTCB for that user thread does not exist.
  - Condition 2: kernel thread is currently assigned
    - ◆ TCB is saved in the memory of the process, AND
    - ◆ KTCB is saved in kernel memory space

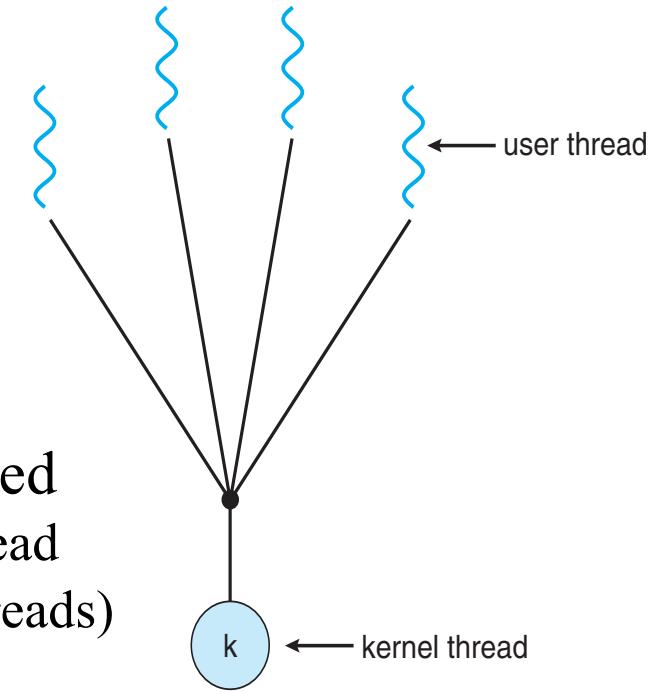
# *Something Subtle Going on Here (2)*

---

- In essence, ALL threads are process threads!
  - Called *user-level threads* to distinguish from kernel threads
- Kernel threads are just containers for user-level threads in order for them to execute
- There can be *MANY* more user-level threads than actual kernel threads ... Why?
- They live “virtually” in the process until they are assigned a kernel thread and execute
  - User-level threads do not take up OS resources!
  - Only kernel threads do (in fact, they ARE the resources)
- Different models are used to assign user-level threads to kernel threads

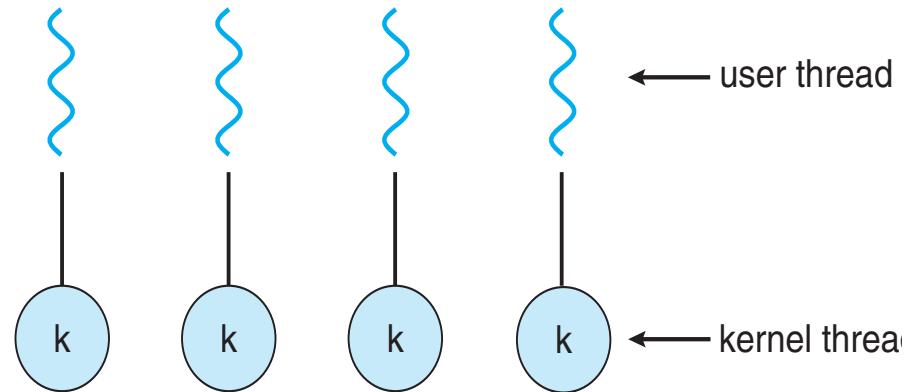
# Many-to-One Thread Model

- Many user-level threads correspond to a *single* kernel thread
  - Kernel/OS is **not** aware of the mapping (OS only sees 1 thread)
  - Handled by a **user-level thread library**
- How does it work?
  - User-level thread library manages:
    - ◆ Create and execute a new thread
    - ◆ Maintain TCBs in process memory (user space)
    - ◆ Ready queue, waiting queues, etc.
    - ◆ Context switch between user-level threads
  - Upon yield, switch to another user thread in the same process
    - ◆ kernel is unaware
  - Upon wait (system calls), all threads are blocked
    - ◆ executing in kernel space by the single kernel thread
    - ◆ kernel is unaware there are other options (user threads)
    - ◆ cannot wait and run at the same time



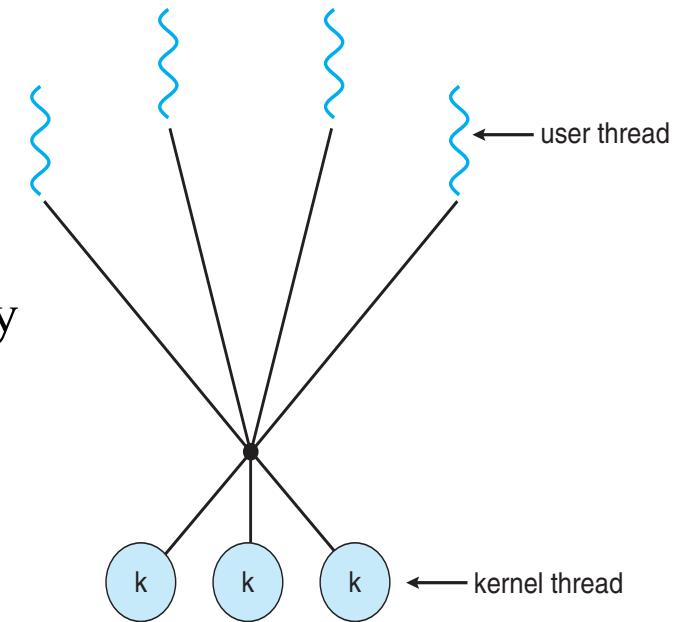
# *One-to-One Thread Model*

- One user-level thread per kernel thread
  - A kernel thread is allocated for **every** user-level thread
  - Must get the kernel to allocate resources for **each** user-level thread
- How does it work?
  - User-level thread library manages:
    - ◆ system call to kernel to create threads
    - ◆ provide user APIs
  - Upon yield, switch to another kernel thread in system
    - ◆ kernel is aware and handle context switch entirely
  - Upon wait, another thread in the process may run
    - ◆ only the single kernel thread is blocked
    - ◆ kernel is aware there are other options in this process



# *Many-to-Many Thread Model*

- A set of user-level threads maps to a set of kernel threads
  - Sizes can be different (kernel set is no larger)
  - No need for the kernel to allocate resources for each new user-level thread
- How does it work?
  - User-level thread library manages:
    - ◆ map user threads to kernel thread dynamically
    - ◆ TCB, queues (same as N:1)
  - Upon yield, switch to another thread
    - ◆ User-level thread library does context switch
  - Upon wait, another thread in the process may run
    - ◆ if a kernel thread is available, user-level can assign a ready user thread to it



# *Thread Assignment*

---

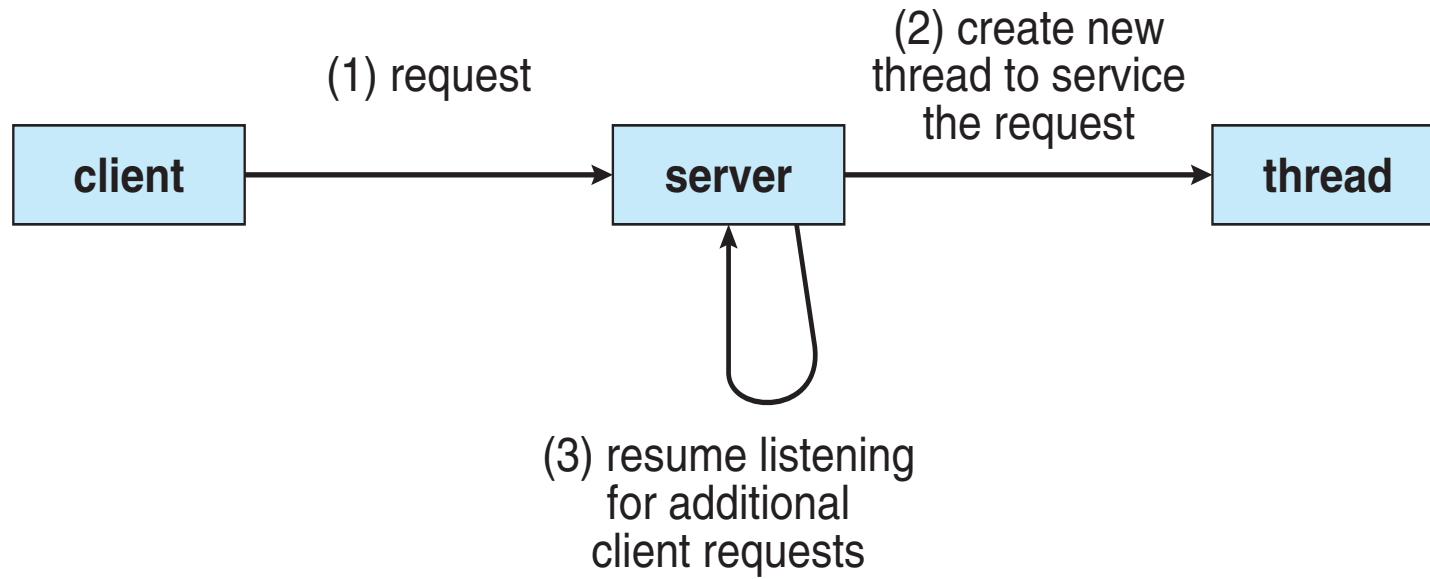
- How many kernel threads should a process have?
  - In an  $M:N$  model there can be significantly more user-level threads ( $M$ ) than kernel-level threads ( $N$ )
  - Kernel threads are the “real” threads that can be allocated to a CPU (core) and run
  - Want to keep enough kernel threads to satisfy the desired level of concurrency in a program and activity in the system
- Suppose that all kernel threads except 1 are blocked
  - What happens if the last kernel thread blocks?
  - Does it matter if there are more user threads “ready” to run?
- In multiprocessing, *thread affinity* is the notion of assigning a thread to run on a particular CPU (core)
  - Help to improve the thread’s performance by keeping its execution resources (CPU, cache, memory) local to CPU

# *Problems Solved with Threads*

---

- Imagine you are building a web server
  - You could allocate a pool of threads, one for each client
    - ◆ thread would wait for a request, get content file, return it
  - How would the different thread models impact this?
- Imagine you are building a web browser
  - You could allocate a pool of threads
    - ◆ some for user interface
    - ◆ some for retrieving content
    - ◆ some for rendering content

# Multithreaded Server Architecture



# Linux Threads

- Linux uses a **one-to-one** thread model
  - Threads are called *tasks*
  - Limit (`cat /proc/sys/kernel/threads_max`)
- Linux views threads as “contexts of execution”
  - Threads are defined separately from processes
  - There is flexibility in what is private and shared
  - So, the line is blur in Linux implements
- Linux system call
  - `clone(int (*fn)(), void **stack, int flags, int argc, ...)`
  - Create a new thread (Linux task)
  - Start running at entry to `(*fn)()`
- May be created in the same address space or not
  - Flags (on means “share”): clone VM, clone filesystem, clone files, clone signal handlers
  - If all these flags off, what system call is clone equal to?

# *POSIX Threads*

- POSIX Threads is a thread API specification
  - Does not define directly the implementation!
  - Could be implemented differently
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
  - Specification, not implementation
  - Provided either as user-level or kernel-level (very interesting)
  - API specifies behavior of the thread library
  - Implementation is up to development of the library
- Common in UNIX operating systems
  - Solaris, Linux, Mac OS X
- POSIX Threads is also known as *Pthreads*



# *POSIX Threads*

---

- *pthread\_create()*
  - start the thread
- *pthread\_self()*
  - return thread ID
- *pthread\_equal()*
  - for comparisons of thread ID's
- *pthread\_exit()*
  - or just return from the start function
- *pthread\_join()*
  - wait for another thread to terminate & retrieve value from *pthread\_exit()*
- *pthread\_cancel()*
  - terminate a thread, by TID
- *pthread\_detach()*
  - thread is immune to join or cancel & runs independently until it terminates
- *pthread\_attr\_init()*
  - thread attribute modifiers

# *POSIX Threads: pthread\_create()*

---

- Interface
- *pthread\_create(...)*
  - Thread ID {of the new thread}
    - ◆ set on return
  - Attributes {of the new thread}
    - ◆ stack size, scheduling information, etc.
  - Function {pointer}
    - ◆ start function for thread
  - Arguments {for function}
  - Return value, status {of the system call}

9

# *POSIX Threads FAQ*

---

- How to pass multiple arguments to start a thread?
  - Build a struct and pass a pointer to it
- Is the pthreads ID unique to the system?
  - No, just process – Linux task ids are system-wide
- After *pthread\_create()*, which thread is running?
  - It acts like fork in that both threads are running
- How many threads terminate when ...
  - *exit()* is called? – all in the process
  - *pthread\_exit()* is called? – only the calling thread
- How are variables shared by threads?
  - Globals, local static, dynamic data (heap)

# *Pthreads Example*

---

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr,"usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr,"%d must be >= 0\n",atoi(argv[1]));
        return -1;
    }
}
```

# Pthreads Example (Cont.)

```
/* get the default attributes */
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid,&attr,runner,argv[1]);
/* wait for the thread to exit */
pthread_join(tid,NULL);

printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

# *Code for Creating/Joining 10 Threads*

```
#define NUM_THREADS 10
/* declare "worker" threads*/
pthread_t workers[NUM_THREADS]

. . .

/* create "worker" threads */
for (int i=0; i<NUM_THREADS; i++)
    pthread_create(&(workers[i]), &attr, runner, argv[1]);

. . .

/* join "worker" threads, one by one */
for (int i=0; i<NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```



See code example  
pthread-simple.c

# *Who executes what?*

- When your program starts
    - OS creates exactly one thread — called the main thread
    - Begins execution at the main() function
  - All other threads created using pthread\_create()
    - Execution at their own thread entry function
    - Not main()
  - When main() function returns
    - Main thread exits
    - By default, the entire process terminates.
    - All other threads are also forcibly terminated, even if they are still running.
    - Exception: call pthread\_exit instead of return in main() → other threads keep running
- ★ See code example  
pthread-main-exit-all.c  
Pthread-main-exit-keep-others.c

# *Meaning of “join” in Threading*

---

- “join” in threading (as in `pthread_join()`), it means:
    - One thread waits for another thread to terminate before continuing.
    - Similar to `wait()` for processes
  - Example: two workers
    - Worker A starts Worker B on a task.
    - Worker A then calls “join” on Worker B.
    - That means A waits until B finishes before A continues.
  - By joining the target thread
    - Clean the target thread’s resources (its stack and internal data)
    - Capture the return value of the target thread
  - Any thread can call `pthread_join()` on another thread
    - Must be created in the same process
    - No other thread has already joined or detached it
- ★ See code example  
`pthread-john.c`

# *Meaning of “detach” in Threading*

---

- When a thread is detached
  - The system will automatically reclaim its resources (stack, control block, etc.)
  - once the thread finishes — you don’t need (and can’t) call `pthread_join()` on it.
- **Joinable thread:** You must later `pthread_join()` it to clean it up.
- **Detached thread:** You cannot join it — it cleans itself up automatically.
- You detach a thread when:
  - You don’t need to know when it finishes.
  - You don’t need its return value.
  - You just want it to “run and go away” (fire-and-forget).
  - Typical for background tasks, logging, cleanup threads, etc.
- A thread can also detach itself!



See code example  
`pthread-detach.c`

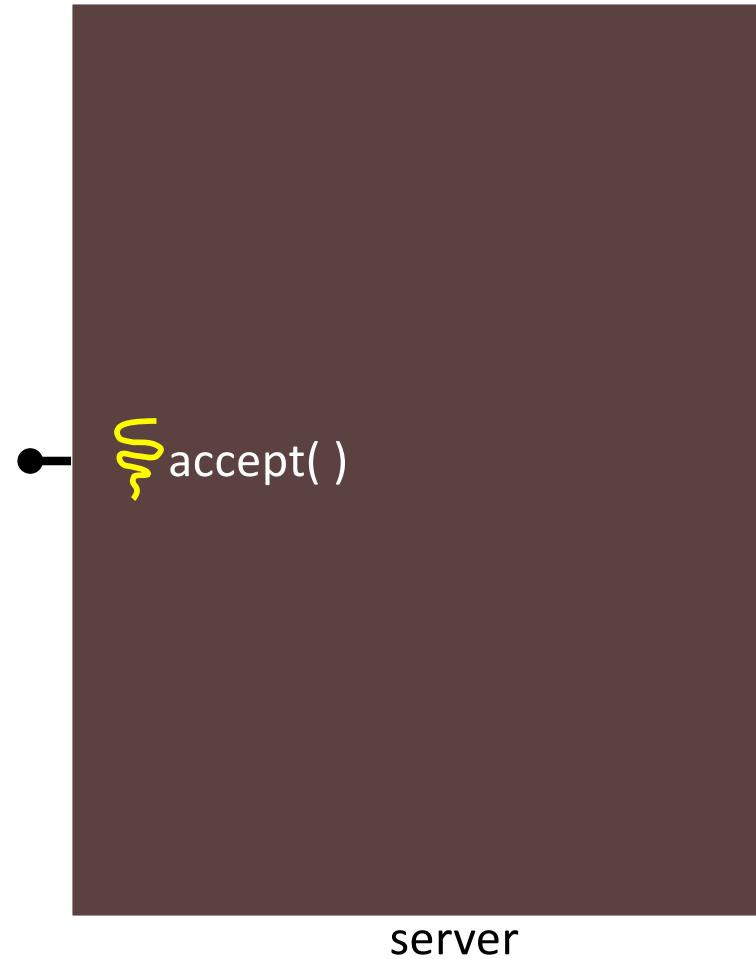
# *Concurrency with Threads*

---

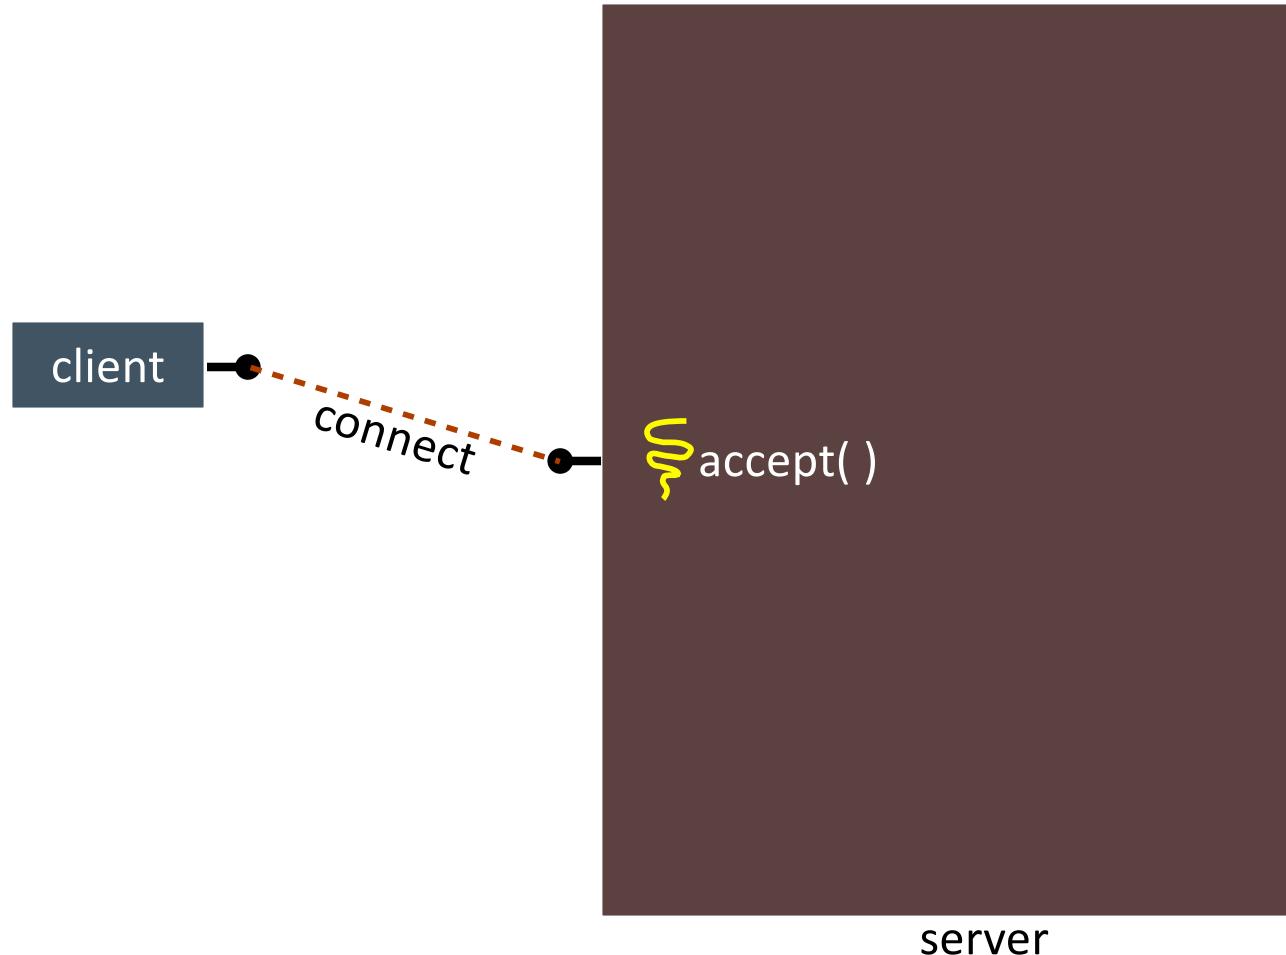
- Consider the client-server example again
- Now want to run with just a single process
  - Need to use threads to get concurrency
- Process “main” (parent) thread waits for clients
  - Parent thread forks (or dispatches) a new thread to handle each new client connection
  - Responsibilities of the child thread:
    - ◆ handles the new connection
    - ◆ exits when the connection terminates

# *Client-Server with Pthreads (1)*

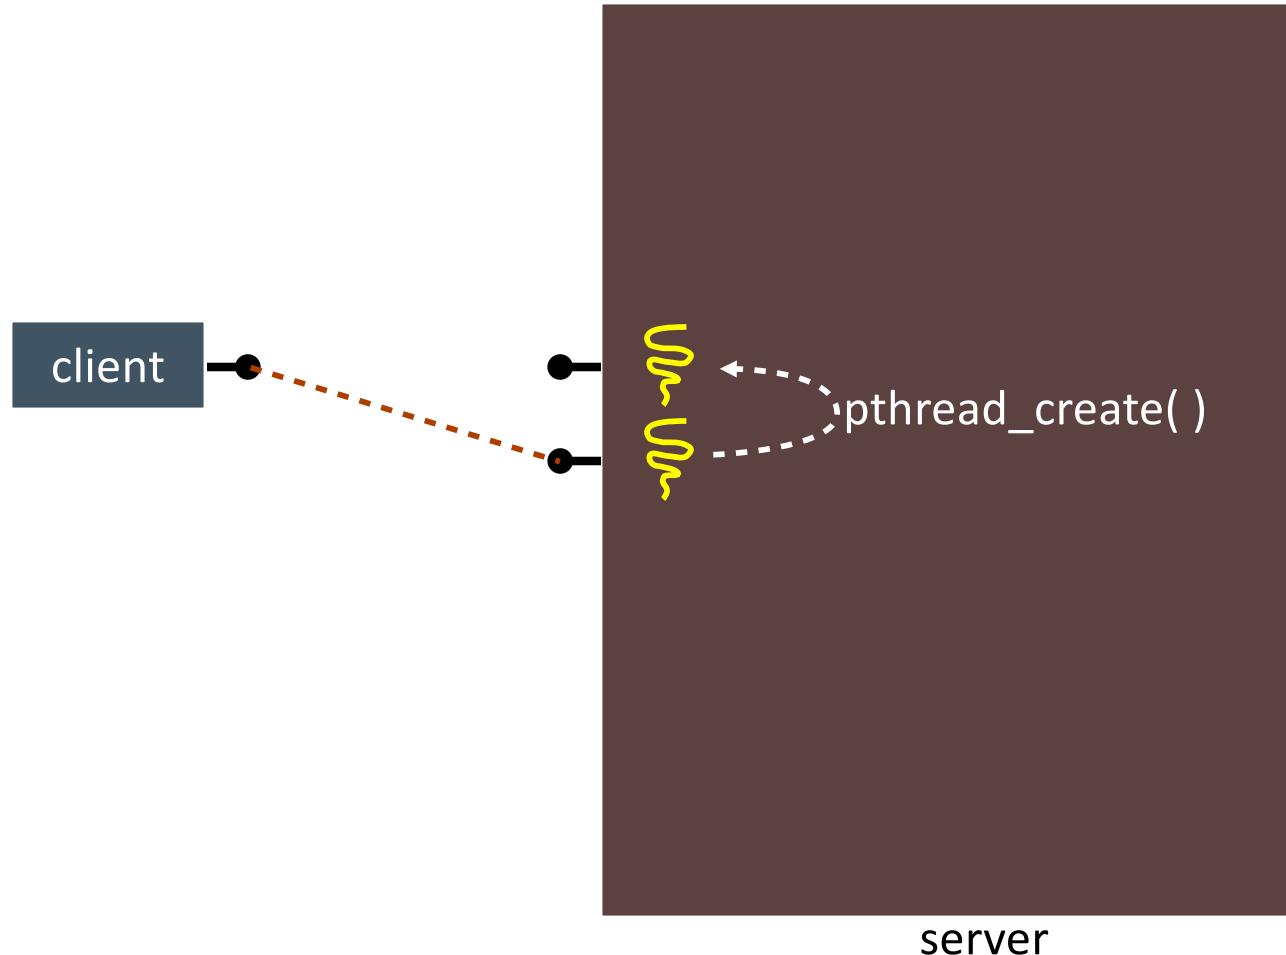
---



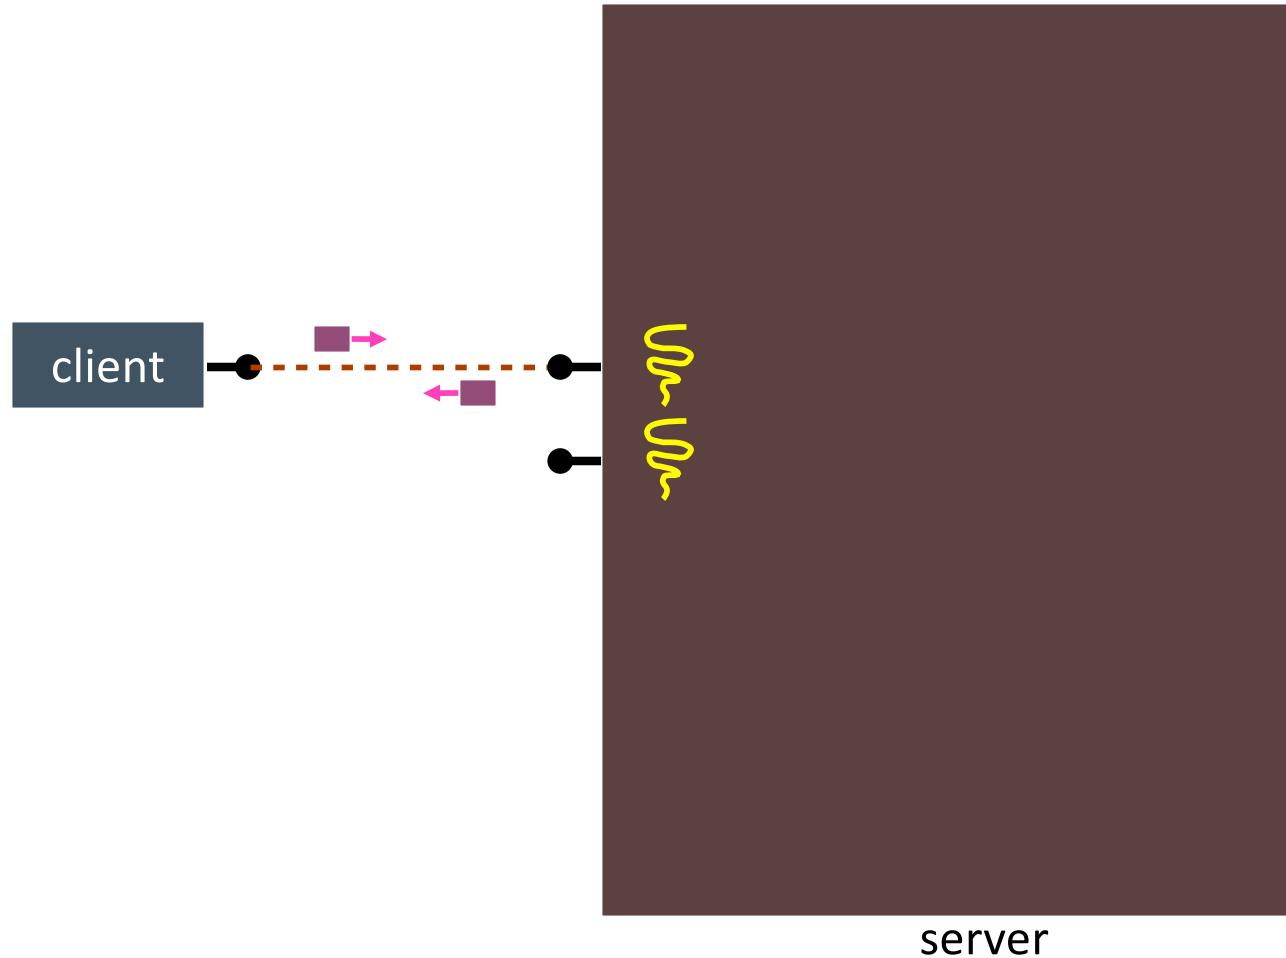
# *Client-Server with Pthreads (2)*



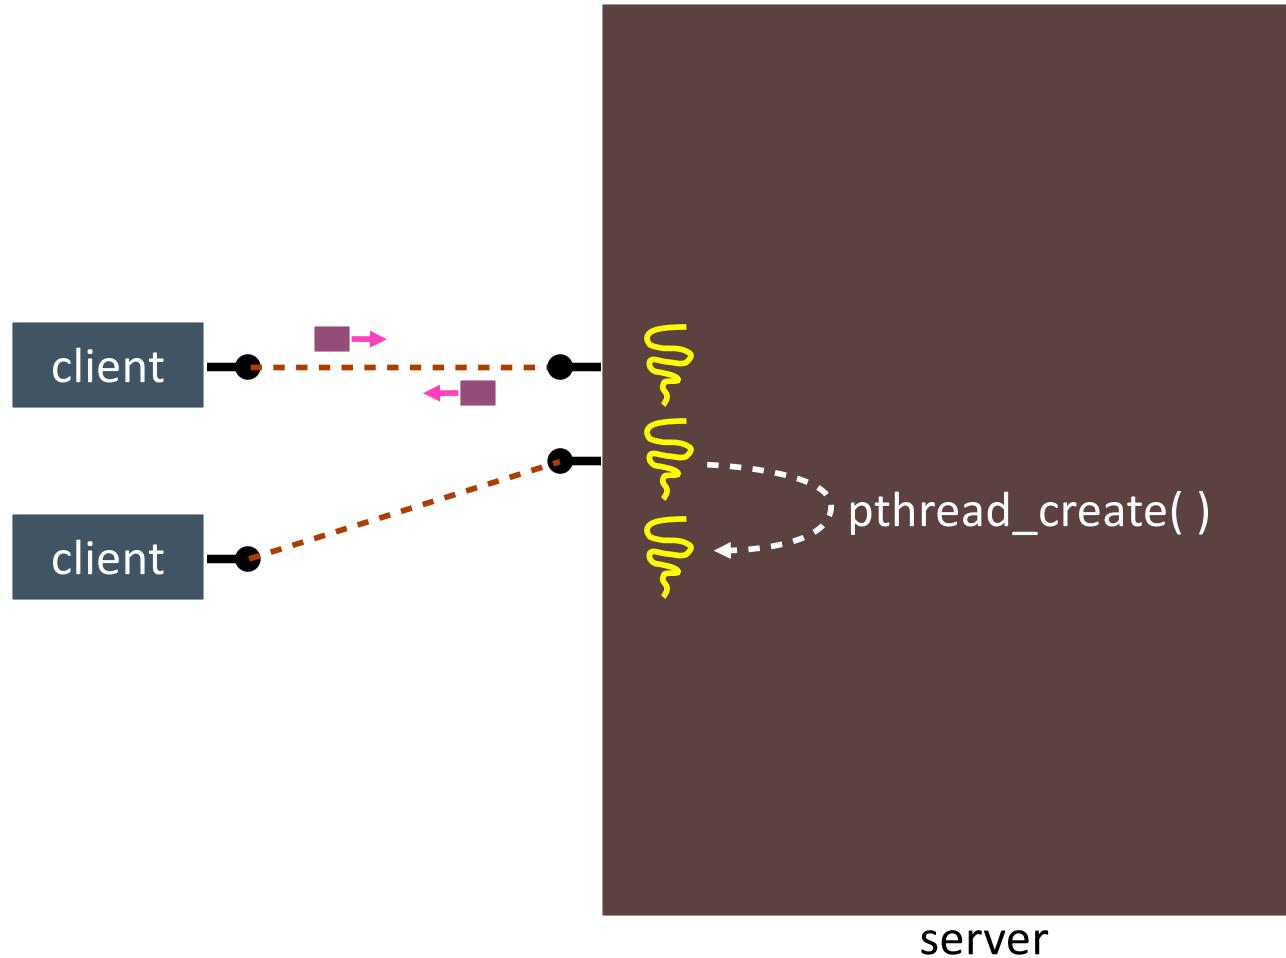
# *Client-Server with Pthreads (3)*



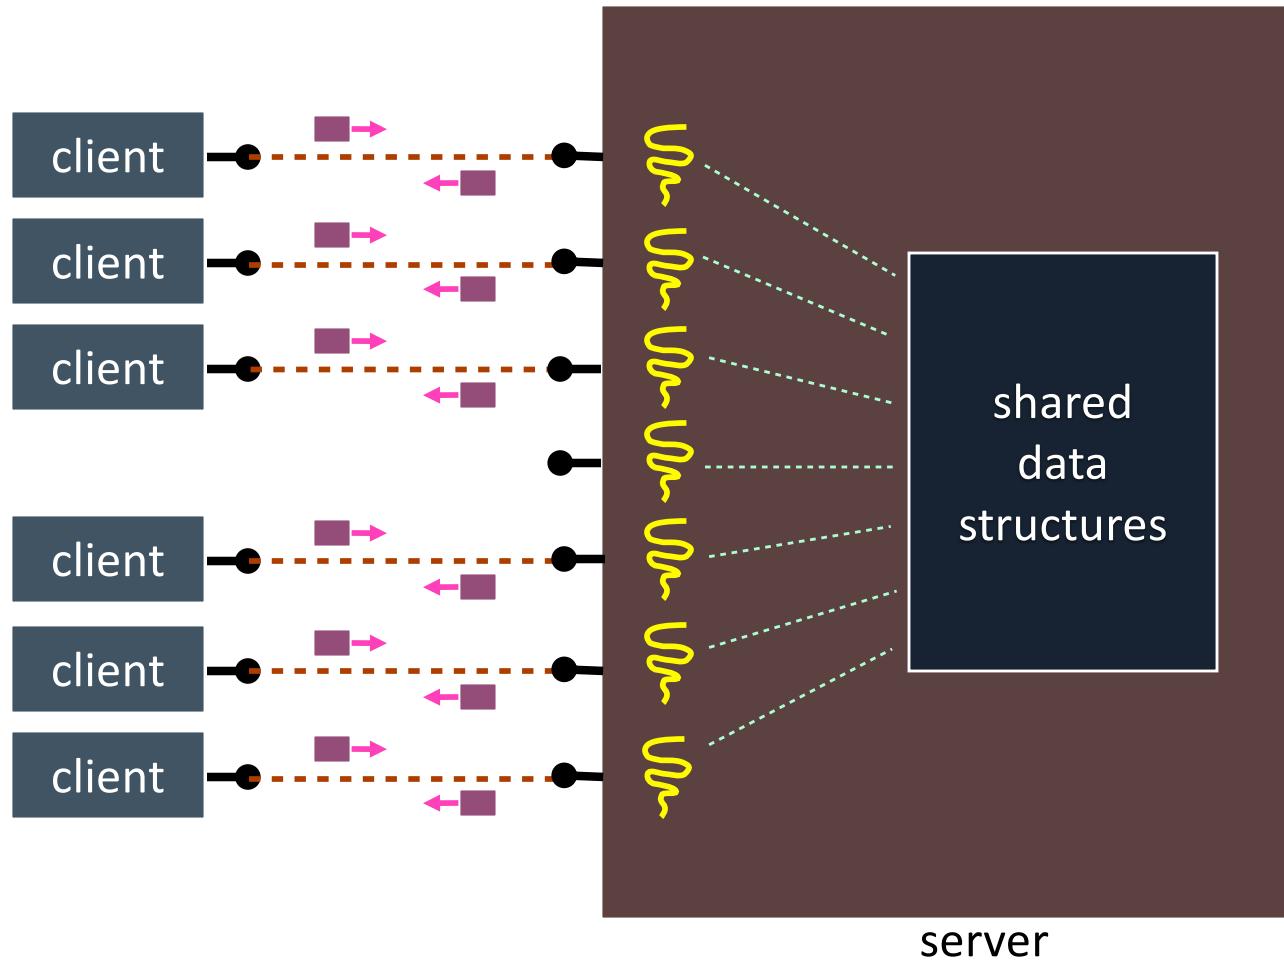
# *Client-Server with Pthreads (4)*



# *Client-Server with Pthreads (5)*



# *Client-Server with Pthreads (5)*



How is this different  
from processes?

# *Implications?*

---

- Consider a web server on a Linux platform
- 0.0297 ms per thread create (time for *clone()*)
  - 10x faster than process forking
  - Maximum of  $(1000 / 0.0297) = \sim 33,670$  connections/sec
  - 3 billion connections per day per machine
    - ◆ much, much better
- Facebook would need only 500 machines
- So, why do we need processes at all?
- Just write everything using threads
  - Writing safe multithreaded code can be complicated
  - Why? What are the issues?

# *Contrasting Processes and Threads*

---

- A process MUST have all necessary information for its execution stored in the OS (process context)
  - Stored in the PCB and takes space in the OS memory
- Each process has a default “thread of execution”
- All processes can execute independently
- A thread is part of (lives in) a process
- All threads in a process share the process context
  - Each thread has a context: PC, stack, CPU state
- All threads are “user” threads in the sense that they have their own PC and a stack
  - This requires very little storage
- For a thread to execute, it must be assigned to a “kernel” thread that will be managed by the OS
  - It can then be context switched (thread switch)

# Compare Overhead of Processes and Threads

- Write a program to create N processes or threads

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

pid_t pid;
int status;

int main(int argc, char *argv[])
{
    int i, numprocesses;

    if(argc != 2) {
        fprintf(stderr, "USAGE: processes <INT>\n");
        exit(1);
    }

    numprocesses = atoi(argv[1]);      // # child processss to create

    for (i=0; i<numprocesses; i++) {
        pid = fork();
        if (pid == 0)           // child process
            exit(0);
        else
            waitpid(pid, &status, 0);
    }
    return 0;
}
```

Processes

```
#include <stdlib.h>
#include <stdio.h>
#include <pthread.h>
#include <sys/time.h>
#include <unistd.h>

struct threadargs {
    int id;
};

struct threadargs      threadARGS;

pthread_t tid;          // thread ID for publishers
pthread_attr_t attr;    // thread attributes
void **retval;

void *routine(void *param);

void *routine(void *args) {
    int tid;           // thread id

    tid = ((struct threadargs *) args)->id;

} // routine

int main(int argc, char *argv[])
{
    int i, numthreads;

    if(argc != 2) {
        fprintf(stderr, "USAGE: threads <INT>\n");
        exit(1);
    }

    numthreads = atoi(argv[1]);      // # child threads to create

    for(i=0; i<numthreads; i++) {
        tid = i;
        pthread_create(&tid, &attr, routine, (void *) &threadARGS);
        pthread_join(tid, retval);
    }
} // main()
```

Threads

- Time for different N

★ See program example

# *Processes versus Threads Performance*

```
ix 95% foreach i (1000 10000 100000 1000000)
foreach? echo Time the creation and shutdown of $i processes
foreach? time processes $i
foreach? echo Time the creation and shutdown of $i threads
foreach? time threads $i
foreach? end

Time the creation and shutdown of 1000 processes
0.209u 0.145s 0:00.33 103.0%    0+0k 46+0io 0pf+0w
Time the creation and shutdown of 1000 threads
0.000u 0.090s 0:00.08 112.5%    0+0k 48+0io 0pf+0w

Time the creation and shutdown of 10000 processes
2.233u 1.452s 0:03.49 105.4%    0+0k 6+0io 0pf+0w
Time the creation and shutdown of 10000 threads
0.086u 0.684s 0:00.67 113.4%    0+0k 8+0io 0pf+0w

Time the creation and shutdown of 100000 processes
22.277u 14.535s 0:35.22 104.4%  0+0k 6+0io 0pf+0w
Time the creation and shutdown of 100000 threads
0.633u 7.107s 0:06.85 112.8%    0+0k 8+0io 0pf+0w

Time the creation and shutdown of 1000000 processes
223.840u 144.968s 5:50.76 105.1%   0+0k 6+0io 0pf+0w
Time the creation and shutdown of 1000000 threads
6.392u 71.093s 1:08.34 113.3%    0+0k 8+0io 0pf+0w
```

N	Processes	Threads
1K	0.33 sec	0.08 sec
10K	3.49 sec	0.67 sec
100K	35.22 sec	6.85 sec
1M	5:50.75 sec	1:08.34 sec

# *Concurrent Threads*

---

- All threads of a process are concurrent
- Benefits
  - All threads are running the same code (of the process)
  - Threads interact directly through shared memory
  - Lower overhead than processes
    - ◆ thread switching is faster
    - ◆ better utilization of the processor, especially with multiple cores
- Disadvantages
  - Maintaining shared memory consistency
    - ◆ synchronization is complicated (see later lecture)
  - One errant thread can affect the whole process
- More to come on this topic ...

# *Thread Pools*

- Pool of threads
  - Create (all threads) at initialization time
  - Assign tasks to run to a waiting thread
    - ◆ it is already made so it should be fast
  - Use all available threads if enough tasks
- What about when a task is done?
  - Should we use terminate the thread executing it?
  - Suppose there is another task to run
  - Should we create another thread for this task?
- Concern for the time to create another thread
  - It is faster than setting up a process
  - But is it necessary?
- How could we improve performance?
  - Just have the current thread wait in a pool



# *Thread Interaction and Communication*

---

- Can you use shared memory?
  - Sure, it is there, might as well use it
  - Just need to allocate memory in the address space
    - ◆ no need for fancy IPC shared memory
- Can you use message passing between threads?
  - Of course, there is nothing to prevent it
  - Would have to build messaging infrastructure to do it
    - ◆ likely would implement in shared memory
    - ◆ not really necessary to use OS mechanisms
- Hmm, can threads utilize IPC mechanisms
  - Certainly, if they were in different processes
  - Would need to make sure only 1 thread uses at a time to avoid consistency problems
  - Strange to use IPC for threads within the same process

# *Thread Cancellation*

---

- So, you want to stop a thread from executing
  - Do not need it anymore
  - It is just hanging around and you want to get rid of it
- Two choices
  - Synchronous (deferred) cancellation
    - ◆ wait for the thread to reach a point where cancellation is permitted
    - ◆ no such operation in Pthreads, but can create your own
  - Asynchronous cancellation
    - ◆ terminate it now
    - ◆ *pthread\_cancel(thread\_id)*

# *Scheduling*

---

- How many kernel threads to create for a process?
  - In M:N model
- If last kernel thread for an application is to be blocked
  - What happens?
  - Recall the relationship between kernel and user threads
- It would be nice if the kernel told the application and the application had a way to get more kernel threads
  - Scheduler activation
    - ◆ at thread block, the kernel tells the application
  - Application can then get a new thread created
    - ◆ see lightweight threads

# *Scheduler Activation*

---

- It would be nice if the kernel told the application that a kernel thread was blocking
  - *Scheduler activation*
    - ◆ at thread block, the kernel tells the application via an *upcall*
    - ◆ an *upcall* is a general term for an invocation of application function from the kernel
  - User-level thread scheduler can then get a new user-level thread created
- Way of conveying information between the kernel and the user-level thread scheduler regarding the disposition of:
  - # user-level threads (increase or decrease)
  - User-level thread state
    - ◆ running to waiting, waiting to ready

# *Fork/Exec Issues*

- Semantics are ambiguous for multithreaded processes
- *fork()*
  - How does it interact with threads?
  - Child process created with only 1 thread?
- *exec()*
  - What happens if process has more than 1 thread?
- *fork* then *exec*
  - Should all threads be copied?



# *Re-entrance and Thread-Safety*

---

- *Re-entrant* code
  - Code that can be run by multiple threads concurrently
  - Can also be safely interrupted and called again (“re-entered”) before the previous execution is finished
- *Thread-safe* libraries
  - Library code that permits multiple threads to invoke the safe function concurrently
  - Mainly concerned with variables that should be private to individual threads
  - Requires moving some global variables to local variables
- Requirements
  - Rely only on input data to the function or data that is already specific to a thread
    - ◆ can not (should not) use any shared data (e.g., global variables) that might be modified by a thread executing the function
  - Must be careful about locking (see later lectures)

# Signal Handling

- What's a signal?
  - A form of IPC
  - Send a particular signal to another process
- Receiver's signal handler processes signal on receipt
- Example
  - Tell the Internet daemon (*inetd*) to reread its config file
  - Send signal to inetd: *kill -SIGHUP <pid>*
  - inetd's signal handler for the SIGHUP signal re-reads the config file
- Note: some signals cannot be handled by the receiving process, so they cause default action (kill the process)



# *Signal Handling*

---

- Synchronous signals
  - Generated by the kernel for the process
  - Due to an exception -- divide by 0
    - ◆ events caused by the thread receiving the signal
- Asynchronous signals
  - Generated by another process
- Asynchronous signals are more difficult for multithreading

# *Signal Handling and Threads*

---

- So, you send a signal to a process
  - Which thread should it be delivered to?
- Choices
  - Thread to which the signal applies
  - Every thread in the process
  - Certain threads in the process
  - A specific signal receiving thread
  - It depends...
- UNIX signal model created decades before Pthreads standard
  - Conflicts arise as a result

# *Why not threads?*

---

- Threads can interfere with one another
  - Need to be careful to maintain consistency
  - Impact of more threads on memory and caches
  - Bug in one thread can cause process problems
- Executing multiple threads may slow them down
  - Impact of single thread vs. switching among threads
- Harder to program a multi-threaded program
  - Multitasking hides context switching
  - Multi-threading introduces concurrency issues

# *Summary of Threads*

---

- Threads
  - A mechanism to improve performance and CPU utilization
- Kernel-space and user-space threads
  - Kernel threads are real, schedulable threads
  - User-space may define its own threads (but not real)
- Threading models and implications
  - Programming systems
  - Multi-threaded design issues
- Threads are useful, but not always the right solution
  - Could slow down system in some cases
  - Can be difficult to program
- Multiprogramming and multithreading are important concepts in modern operating systems

# *Next Class*

---

## □ Scheduling