



CS 415

Operating Systems

Processes

Prof. Jieyang Chen

Department of Computer and Information Science

Fall 2025



UNIVERSITY OF OREGON

Logistics

- Project 1 posted
- Labs
 - Intended to support projects
 - Lab 2 posted on Monday
- Read OSC Chapter 3
- Make use of Canvas discussion
- Get familiar with the Linux man pages
<https://man7.org/linux/man-pages/index.html>

Outline

- Process concept
- Process operation
- System calls to create processes
- Process management
- Process scheduling

Overview of Processes

- We have programs, so why do we need processes?
- Questions that we explore
 - How are processes created?
 - ◆ from binary program to executing process
 - How is a process represented and managed?
 - ◆ process creation, process control block
 - How does the OS manage multiple processes?
 - ◆ process state, ownership, scheduling
 - How can processes communicate?
 - ◆ interprocess communication, concurrency, deadlock

Superview and User Modes

- OS runs in “supervisor” mode
 - Has access to protected (privileged) instructions only available in that mode (kernel executes in ring 0)
 - Allows it to manage the entire system
- OS “loads” programs into processes
 - User programs run in user mode
 - Many processes can run in user mode at the same time
- How does OS get programs loaded into processes in user mode and keep them straight?

Process Concept

- A process is a *program in execution*
 - Process execution can result in more processes being created
- What makes up a process?
 - Program executable code (called *image* or *text*)
 - CPU state (program counter, CPU registers, ...)
 - Stack containing temporary data (as a result of function calls)
 - ◆ function parameters, return addresses, local variables, ...
 - ◆ called *stack frames*
 - Data section containing global variables
 - Heap containing memory dynamically allocated during run time

Process Concept

- A process is a *program in execution*
 - Process execution can result in more processes being created
- What makes up a process?
 - Code –instructions needed to execute
 - Data – data needed to complete the task
 - Execution context – what we are executing right now
 - Management information – what files are opened, what IO devices are used, etc.

All together is called **process context**


Process context is everything a process needs

Process Context

- *How is process context implemented? Where is it?*
 - Code + Data: **Process address space** (in memory)
 - Execution context: **CPU state** (program counter, CPU registers)
 - Management information: **Process Control Blocks** (in memory)
- **Process address space**: code, stack, data, heap
- **CPU states**
 - Updates only when executing
 - When not currently executing, it must be saved somewhere
 - (will go back to this when talking about context switch)
- **Process Control Blocks**
 - Process ID and state
 - Memory management info
 - Info about current usage on disk, files, ...
 - Scheduling info: priority


Process Address Space

- *Process address space* is all locations addressable by the process
 - Each process has its own address space
- Code (*Text*)
- Global Data (*Data*)
- Dynamic Data (*Heap*)
 - Grows up
- Local Data (*Stack*)
 - Grows down



Process Address Space

- A process has to reference memory for different purposes
 - Getting instructions (via PC)
 - Access data: static/global (data) and dynamic (heap)
- Where are these things located?
- Logical address
 - # address bits in instructions determine logical memory size
 - 48-bit can address up to 2^{48} bytes
 - A “logical address” is from 0 to the size of logical memory
 - Compiler and OS determine where things get placed in logical memory (will be back to this when we talk about virtual memory)



Process Address Space

```
int value = 5; → Global/Data
```

```
int main()
{
    int *p; → Stack
```

```
p = (int *)malloc(sizeof(int)); → Heap
```

```
if (p == 0) {
    printf("ERROR: Out of memory\n");
    return 1;
}
```

```
*p = value;
printf("%d\n", *p);
free(p);
return 0;
}
```

Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...){
    int nums[4] = {2,4,6,8};
    int *ncopy = copy(nums, 4);
    // ... do stuff ...
    free(ncopy);
    return 0;
}
```



Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...){
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(ncpy);
    return 0;
}
```



Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...){
    int nums[4] = {2,4,6,8};
    int *ncopy = copy(nums, 4);
    // ... do stuff ...
    free(ncopy);
    return 0;
}
```



Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...){
    int nums[4] = {2,4,6,8};
    int *ncopy = copy(nums, 4);
    // ... do stuff ...
    free(ncopy);
    return 0;
}
```



Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...){
    int nums[4] = {2,4,6,8};
    int *ncopy = copy(nums, 4);
    // ... do stuff ...
    free(ncopy);
    return 0;
}
```



Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...){
    int nums[4] = {2,4,6,8};
    int *ncopy = copy(nums, 4);
    // ... do stuff ...
    free(ncopy);
    return 0;
}
```



Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...){
    int nums[4] = {2,4,6,8};
    int *ncopy = copy(nums, 4);
    // ... do stuff ...
    free(ncopy);
    return 0;
}
```



Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...){
    int nums[4] = {2,4,6,8};
    int *ncopy = copy(nums, 4);
    // ... do stuff ...
    free(ncopy);
    return 0;
}
```



Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(nncpy);
    return 0;
}
```



Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(... {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(nncpy);
    return 0;
}
```



Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2, 4, 6, 8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(nncpy);
    return 0;
}
```



Heap + Stack


```
#include <stdlib.h>

int *copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(
        size * sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];
    return a2;
}

int main(...) {
    int nums[4] = {2,4,6,8};
    int *ncpy = copy(nums, 4);
    // ... do stuff ...
    free(nncpy);
    return 0;
}
```



Summary on Process Addresses

- A process executes instructions
- Addresses generated by the instructions
 - They could reference data
 - They could reference the stack
 - They could reference another instruction
- Program counter is an instruction address
- Stack pointer (stack register) is a stack address
- All addresses produced by process instruction execution on a CPU are **logical addresses**
- Actual physical addresses to physical memory are **translated** from logical address
 - Done by *memory address translation hardware*

Program to Process

- Program is stored in a binary format
 - Executable and Linkable Format (ELF)
 - a.out
- Binary format describes
 - Program sections: Text, Data, ... (many types of sections)
 - What to load at execution time
- Determined by the compiler, loader, linker, and OS

ELF Files

+-----+ ELF Header ← Identifies ELF file type, architecture, entry point +-----+
+-----+ Program Header Table ← Describes loadable segments for the loader +-----+
+-----+ Section Contents ← Various sections used for linking/debugging .text Machine code (instructions) .rodata Read-only constants .data Initialized global/static variables .bss Uninitialized global/static variables More... +-----+
+-----+ Section Header Table ← Metadata describing each section +-----+

Why is there no stack/heap section?

- ELF sections describe **static content of the program** (what is known **before execution**)
- Their size and usage are **determined at runtime**.
- Created **dynamically by the OS when the program starts**.

CPU states

What is required to execute each instruction on CPU?

- Registers store the state of execution in CPU
 - Program counter (there is more actually...)
 - Data registers
- Program counter to indicate what to execute?
 - Holds address/index of next instruction to execute

```
int main() {  
    int a = 5;  
    int b = 3;  
    int c = a + b;  
    return c;  
}
```

C program

```
Addr 100: MOV R1, 5 ; Load a = 5  
Addr 101: MOV R2, 3 ; Load b = 3  
Addr 102: ADD R3, R1, R2 ; c = a + b  
Addr 103: MOV R0, R3 ; return value in R0  
Addr 104: HALT
```

Compile to assembly (simplified)

CPU states

```
Addr 100: MOV R1, 5 ; Load a = 5
Addr 101: MOV R2, 3 ; Load b = 3
Addr 102: ADD R3, R1, R2 ; c = a + b
Addr 103: MOV R0, R3 ; return value in R0
Addr 104: HALT
```

Compile to assembly (simplified)

CPU states



```
Addr 100: MOV R1, 5 ; Load a = 5
Addr 101: MOV R2, 3 ; Load b = 3
Addr 102: ADD R3, R1, R2 ; c = a + b
Addr 103: MOV R0, R3 ; return value in R0
Addr 104: HALT
```

Compile to assembly (simplified)

CPU state:

PC = 100

R0 = ?

R1 = ?

R2 = ?

R3 = ?

CPU states



```
Addr 100: MOV R1, 5 ; Load a = 5
Addr 101: MOV R2, 3 ; Load b = 3
Addr 102: ADD R3, R1, R2 ; c = a + b
Addr 103: MOV R0, R3 ; return value in R0
Addr 104: HALT
```

Compile to assembly (simplified)

CPU state:

PC = 101

R0 = ?

R1 = 5

R2 = ?

R3 = ?

CPU states



```
Addr 100: MOV R1, 5 ; Load a = 5
Addr 101: MOV R2, 3 ; Load b = 3
Addr 102: ADD R3, R1, R2 ; c = a + b
Addr 103: MOV R0, R3 ; return value in R0
Addr 104: HALT
```

Compile to assembly (simplified)

CPU state:

PC = 102

R0 = ?

R1 = 5

R2 = 3

R3 = ?

CPU states



```
Addr 100: MOV R1, 5 ; Load a = 5
Addr 101: MOV R2, 3 ; Load b = 3
Addr 102: ADD R3, R1, R2 ; c = a + b
Addr 103: MOV R0, R3 ; return value in R0
Addr 104: HALT
```

Compile to assembly (simplified)

CPU state:

PC = 103

R0 = ?

R1 = 5

R2 = 3

R3 = 8

CPU states

```
Addr 100: MOV R1, 5 ; Load a = 5
Addr 101: MOV R2, 3 ; Load b = 3
Addr 102: ADD R3, R1, R2 ; c = a + b
Addr 103: MOV R0, R3 ; return value in R0
Addr 104: HALT
```



Compile to assembly (simplified)

CPU state:

PC = 104

R0 = 8


R1 = 5

R2 = 3

R3 = 8

Process State


- As a process executes, it changes state
 - *New*: The process is being created (starting state)
 - *Running*: Instructions are being executed
 - *Waiting*: The process is waiting for some event to occur
 - *Ready*: The process is waiting to run
 - *Terminated*: The process has finished execution (end state)



Process Control Block


- Information associated with each process
 - Also called the *task control block*
- Process state: running, waiting, ...
- Program counter
 - Location of instruction to execute next
- CPU registers for process thread (data)
- CPU scheduling information
 - Priorities, scheduling queue pointers, ...
- Memory-management information
 - Memory allocated to the process
- Accounting information
 - CPU used, clock time elapsed, ...
- I/O status information
 - I/O devices allocated to process
 - List of open files

PCB



Scheduling Processes


- Processes transition among execution states



Certain operating systems also more finely delineate process states


State Transitions

- New Process ==> Ready
 - Allocate resources necessary to run
 - Place process on ready queue (usually at end)
- Ready ==> Running
 - Process is at the head of ready queue
 - Process is scheduled onto an available processor
- Running ==> Ready
 - Process is interrupted
 - ◆ usually by a timer interrupt
 - Process could still run, in that it is not waiting something
 - Placed back on the ready queue



State Transitions (2)

- Running ==> Waiting
 - Either something exceptional happened that caused an interrupt to occur (e.g., page fault exception) ...
 - ... or the process needs to wait on some action (e.g., it made a system call or requested I/O)
 - Process must wait for whatever event happened to be serviced
- Waiting ==> Ready
 - Event has been satisfied so that the process can return to run
 - Put it on the ready queue
- Ready ==> Running
 - As before...



Next Class

- Process scheduling