

Memory

<p>Swapping</p> <p>What is Swapping? Temporarily moves process (or part) to backing store and reloads to main mem as needed.</p> <ul style="list-style-type: none"> Allows total address space > physical mem → increases multiprogramming. With process scheduling (*roll out/roll in* lower priority processes). <p>Why use Swapping?</p> <ul style="list-style-type: none"> Run programs larger than physical mem. Increase degree of multiprogramming. Improve CPU utilization & throughput (more ready processes). Efficient process mem allocation. Reduce I/O vs loading full program at once. <p>Standard Swapping Moves entire process image between main mem & backing store.</p> <p>Disadvantages:</p> <ul style="list-style-type: none"> High transfer time for large processes (included in context switch time). <p>Swap Space Management</p> <p>Secondary storage for swapping pages. Raw partition or large file. Primarily backs stack/heap. Executable pages → demand-paged from file system & discarded if unmodified.</p> <p>Swapping with Paging (Demand Paging) Swaps individual pages instead of full process.</p> <ul style="list-style-type: none"> If page not in mem → page fault → load page from backing store. No free frames → page-replacement selects victim frame. Modified victim written back; unmodified discarded. <p>Pros: Run large virtual address space on small physical mem. Cons: Frequent page faults → performance hit.</p>	<p>Paging</p> <p>Purpose: Allows non-contiguous physical address space ↳ solves external fragmentation.</p> <p>Key Idea: Logical page can reside in any physical frame.</p> <p>Logical Address Space → Pages</p> <p>Physical Mem → Frames Pages loaded into any available frame.</p> <p>Steps of MMU:</p> <ol style="list-style-type: none"> Extract page number (p) from logical address. Index into page table → frame number (f). MMU forms physical address → frame f + page offset (d). <p>Page Table:</p> <ul style="list-style-type: none"> Maps logical page numbers → physical frame numbers. Stored in main mem → PTBR points to it. Each entry: frame number + protection bits + valid-invalid bit. <p>Valid-Invalid Bit:</p> <ul style="list-style-type: none"> Valid: page in mem. Invalid: page not in mem → triggers page fault. <p>Paging Hardware & TLB</p> <ul style="list-style-type: none"> Without TLB → 2 mem accesses per reference (page table + data). TLB: Cache of recent page # → frame # translations. <p>TLB hit: quick access. TLB miss: must access page table → load translation into TLB. ASID: lets TLB entries coexist for multiple processes.</p>	<p>Fragmentation Fragmentation → Wasted space due to imperfect fit b/w allocated space and requested space</p> <p>Two types:</p> <ul style="list-style-type: none"> External Fragmentation <ul style="list-style-type: none"> Free space exists but not contiguous → can't satisfy request. Occurs in contiguous allocation, segmentation, contiguous file allocation. Example: lots of small holes scattered in memory. Internal Fragmentation <ul style="list-style-type: none"> Allocated block > requested size → unused space inside allocated block. Happens with paging (last frame), fixed-size allocations, clustered file allocation, buddy system. <p>Where it appears:</p> <ul style="list-style-type: none"> Main Memory <ul style="list-style-type: none"> Contiguous allocation → external fragmentation (small scattered holes). Paging → internal fragmentation (last frame of process). Buddy system → internal fragmentation from rounding to power of 2. File Systems <ul style="list-style-type: none"> Contiguous allocation → external fragmentation. Linked allocation, indexed allocation → no external fragmentation, but internal if using large clusters. Extents reduce external frag but large extents can still waste space. <p>Solutions:</p> <ul style="list-style-type: none"> External frag → compaction (expensive, requires dynamic relocation). Internal frag → tune page/block/cluster size to balance overhead vs. fragmentation. <p>File System Allocation Methods</p> <p>Contiguous Allocation: Each file = set of contiguous blocks. Pros: fast sequential access, easy direct access. Cons: external fragmentation; hard to grow file.</p> <p>Linked Allocation: File → linked list of blocks. Pros: no external fragmentation. Cons: poor random access; pointer overhead.</p> <p>Indexed Allocation: File → index block with pointers to data blocks. Pros: efficient direct access, no external fragmentation. Cons: index block overhead.</p> <p>Large Files:</p> <ul style="list-style-type: none"> Linked Index Blocks Multilevel Index Combined Scheme (UNIX Inode): direct, single, double, triple indirect pointers. <p>Thrashing Cause: Too many processes → excessive paging. Cycle: • New process faults → steals frames → more faults → CPU idle → OS adds more → worse. Detection: High page fault rate + low CPU utilization. Solutions: <ul style="list-style-type: none"> Decrease degree of multiprogramming. Local replacement → isolate offending process. Use Working Set Model or Page Fault Frequency (PFF) to manage frame allocation. </p>
<p>Allocation Methods</p> <p>Contiguous Mem Allocation Each process occupies single block of physical mem.</p> <ul style="list-style-type: none"> Uses base & limit registers → maps logical to physical addresses → enforces protection. <p>Allocation Strategies (Dynamic Storage Allocation)</p> <ul style="list-style-type: none"> First Fit: First hole large enough (faster than Best Fit). Best Fit: Smallest hole large enough; smallest leftover hole. Worst Fit: Largest hole; largest leftover hole. First/Best > Worst for speed & utilization. <p>Non-Contiguous Mem Allocation</p> <ul style="list-style-type: none"> Segmentation: Program divided into logical segments → segment table: base address + limit. Logical address = <segment #, offset>. Suffers external fragmentation. Paging: Physical mem → fixed-size frames. Logical mem → fixed-size pages. No external fragmentation. <p>Effective Access Time (EAT)</p> <p>Formula: $EAT = (\text{hit ratio} * (\text{TLB access} + \text{mem access})) + (\text{miss ratio} * (\text{TLB access} + \text{page table access} + \text{mem access}))$</p> <p>Impact: High TLB hit ratio → minimizes extra overhead. Without TLB, paging doubles access time.</p>	<p>Virtual Memory</p> <p>Definition: Execution of processes not fully in mem. Separates logical from physical mem.</p> <p>Benefits:</p> <ul style="list-style-type: none"> Run larger programs than physical mem. More programs concurrently. Less I/O (unused pages not loaded). Shared libs & files → page sharing. Efficient process creation (copy-on-write). <p>Mem Initialization: Zero-out free frame → avoid security issues.</p> <p>Page Fault Handling</p> <p>Detection: MMU sees invalid bit → triggers trap.</p> <p>Handling Steps:</p> <ol style="list-style-type: none"> Save registers/process state. Validate page ref → locate page on disk. Find free frame → schedule I/O. Load page → update page table → set valid bit. Resume/restart instruction. <p>No Free Frame: page replacement required. If victim is dirty → write back first → adds delay.</p> <p>Performance Impact: Page faults → millisecond delays (vs nanosecond mem access). High fault rate → poor performance.</p>	<p>Working Set</p> <p>Working Set Window (Δ): defines current working set.</p> <p>WSS (Working Set Size): number of pages referenced in Δ.</p> <p>If $\Sigma WSS_i > \text{total frames}$: thrashing occurs → reduce active processes.</p> <p>Tracking: difficult → approximated w/ ref bits + timer interrupts.</p> <p>Behavior: process entering new locality → temporary page fault spike → settles when working set loaded.</p> <p>Using a hierarchical page table helps to solve what problems?</p> <ul style="list-style-type: none"> X Thrashing: Not directly solved by page table structure; thrashing is caused by insufficient memory frames. Large page table sizes: Correct — hierarchical page tables break large tables into manageable parts. X Belady's anomaly: Caused by certain page replacement algorithms, unrelated to page table structure. X Memory sharing: Achieved by mapping multiple processes to same frames, not by hierarchy. X External fragmentation: Paging already eliminates external fragmentation. <p>Which steps are NOT involved in servicing a page fault?</p> <ul style="list-style-type: none"> Flush the TLB: Incorrect (This IS a step; old translation must be flushed after loading the page.) Find a free frame: Incorrect (IS a step; must load page into a frame.) Set dirty bit to TRUE: Correct — dirty bit is set only when writing happens later, not on page fault. Terminate process: Correct — page fault is normally handled by loading the page, not by killing process. Update page table entry: Incorrect (IS a step; must mark page as present.) Write referenced page to backing store: Correct — backing store is written only if evicting a dirty page, not as part of initial fault servicing. <p>Which of the following could be beneficial to improving the efficiency of a virtual memory system?</p> <ul style="list-style-type: none"> Smaller page sizes: Incorrect Not necessarily; too small pages increase page table size and overhead. Larger TLB: Correct — larger TLB improves hit ratio, reduces memory access overhead. Larger backing store: Incorrect Larger backing store capacity does not improve speed; faster is what matters. Faster backing store: Correct — reduces page fault service time. Page prefetching: Correct — prefetching reduces page fault frequency. Fewer active processes: Correct — fewer processes reduce memory pressure and page faults.
<p>Memory Address Translation</p> <p>Purpose: Maps logical → physical addresses. Enables multiple processes in mem concurrently & ensures mem protection.</p> <p>Logical vs Physical Addresses</p> <ul style="list-style-type: none"> Logical Address (Virtual Address) — generated by CPU; part of process's <i>logical address space</i>. Physical Address — actual location in main mem; part of <i>physical address space</i>. <p>Address Binding</p> <ul style="list-style-type: none"> Compile Time — If mem location known, generate absolute code. If moved → recompile. Load Time — If location unknown, generate relocatable code. Binding happens at load time. Execution Time — Binding done during execution; allows moving process in mem. Requires hardware support. <p>Protection & Relocation (Base + Limit Registers)</p> <ul style="list-style-type: none"> CPU checks each address against base and limit registers → prevents illegal access. Registers only modifiable by OS; MMU adds base to logical address → forms physical address. <p>Although segments are a non-contiguous memory management approach, external fragmentation is still possible.</p> <p>True. External fragmentation is still possible in segmentation because free memory is divided into variable-sized segments, which can leave small unusable gaps between allocated segments.</p> <p>The translation lookaside buffer (TLB) is stored in main memory.</p> <p>False The TLB is a hardware cache stored inside the CPU, not in main memory.</p> <p>The # entries in an inverted page table is equal to the number of frames in logical memory.</p> <p>False The number of entries in an inverted page table equals the number of physical frames, not logical memory size.</p> <p>The amount of virtual memory used by a process can exceed physical memory.</p> <p>True Virtual memory allows a process to use more memory than is physically available by swapping pages in and out of disk.</p> <p>The working set model is an approach used to help manage the number of frames a process needs in order to execute smoothly and avoid thrashing.</p> <p>True The working set model tracks the set of pages a process actively uses to ensure enough frames are allocated to prevent thrashing.</p>	<p>Page Table Structure</p> <p>For large address spaces, page tables can become very large. Fixes:</p> <ul style="list-style-type: none"> Hierarchical Paging (Multi Level) <ul style="list-style-type: none"> Dividing page table into smaller tables, often by paging page table itself, addresses are translated by traversing multiple levels of page tables Hashed Page Tables <ul style="list-style-type: none"> Hash function maps page #s to hash table entries, which point to page table entries; useful for sparse address spaces Inverted Page Tables <ul style="list-style-type: none"> Having one entry per physical frame, mapping virtual addrs and PID stores in frame. Decreases mem for page tables but increases lookup time (often uses hash table → faster search) <p>Free Space Management</p> <p>OS needs to manage list of free blocks (or clusters) on storage device</p> <ul style="list-style-type: none"> Bit Vector (Bitmap) <ul style="list-style-type: none"> Each block represented by bit (1 free, 0 allocated), efficient for finding first free block or contiguous free blocks, requires extra space for bitmap Linked List (Free List) <ul style="list-style-type: none"> All free blocks linked together with pointers, first free block points to next, so on, no wasted space for a map, not efficient for finding contiguous space Grouping <ul style="list-style-type: none"> Modification of linked list where first free block stores addys of next n free blocks, with last pointer pointing to next block containing addresses, improves speed of finding free blocks Counting <ul style="list-style-type: none"> Takes advantage of contiguous free blocks by storing addys of first free block and count of contiguous blocks that follow it. Each entry is an address/count pair, more efficient than a simple linked list if counts > 1, similar extent to allocation Space Maps <ul style="list-style-type: none"> Divides space into chunks and uses space maps per chunk. Uses counting format and log structured techniques for updates. Designed for very large file systems <p>Page Replacement</p> <p>Goal: Select victim page → minimize future page faults.</p> <p>Belady's Anomaly: More frames can → more faults.</p> <p>Algorithms:</p> <ul style="list-style-type: none"> Optimal (OPT/MIN): Replace page not needed for longest time (theoretical). FIFO: Oldest page replaced → can evict frequently used pages. LRU Approximation: <ul style="list-style-type: none"> Second Chance (Clock): Pages w/ reference bit set → given second chance. Enhanced Second Chance: Uses both ref bit + modify bit → 4 classes. 	