

# Section 2 : Scheduling

## When does scheduling happen?

The OS makes a scheduling decision when one of these things happen :

- Running → Waiting (ie process makes an I/O request)
- Running → Ready (timer interrupt in Round Robin)
- Waiting → Ready (I/O finishes)
- Process terminates
- New process created

## Scheduling Goals (What are we trying to optimize?)

- Keep the CPU busy 100% of the time
- Throughput - # of processes completed per time unit
- Turnaround time - time from submission to completion
- Waiting time - Time from submission to first CPU execution
- Response time - Time from submission to first CPU execution
- Fairness - No processor starves, all get a turn, everyone eats

## Resource types

### Memory

#### Physical memory (RAM)

- Limited, only so much space for programs to live in
- OS allocates memory when a program starts and frees it when it exits

#### Virtual memory

- Tricks programs into thinking that they have more mem than is physically available
- Allows multiple processes to co-exist by using disk as backup
- Enable memory protection and process isolation

#### I/O - disk drives, printers network cards

- + :: I/O is slow so it must be scheduled carefully, can request I/O and block until its done

#### CPU

- Time shared, only one process can run per core at a time
- OS uses scheduling algorithms to decide who gets the CPU next

## Multilevel Queue (MLQ)

**How it works:** Multiple separate queues for different process types

- i.e. System processes, interactive, batch
- Each queue can use its own scheduling algorithm
- Queue scheduling order is fixed (ie system → interactive → batch)

**Pros:** Great for managing process types with different needs

**Cons:** Rigid - processes are stuck in one queue

- Starvation possible if low priority queues are never served

**Def : Starvation** - When a process waits indefinitely to get the CPU because other processes keep getting chosen instead

## Multilevel Feedback Queue (MLFQ)

**How it works:** Like MLQ, but processes can move between queues

- Processes that use too much CPU time get "demoted"
- New/interactive processes start in higher-priority queues
- More flexible and adaptive

**Pros:** Adjust priority based on behavior, Very responsive for short jobs and interactive processes

**Cons:** Hard to configure, Can be gamed if not implemented carefully

- Starvation is possible, but is mitigated through policy

## CPU Scheduling Algorithms

### First come first serve (FCFS)

**How it works:** Processes run in the order they arrive

- No preemption - each process runs to completion

**Pros:** Simple to implement, Fair in order of arrival

**Cons:** Convoy effect - One long job delays everything, Poor responsiveness for short processes.

- No starvation - every process will eventually get the CPU

### Shortest Job First (SJF)

**How it works :** Pick the process with the shortest estimated CPU burst

Can be either Non preemptive (Classic), Preemptive (Shortest Remain Time First SRTF)

**Pros:** Optimal average waiting time

**Cons:** Need to predict burst times, hard, Starvation of long job if short jobs keep arriving

### Round Robin (RR)

**How it works :**

- Fixed time quantum for each, If it doesn't finish in time, sent to back of the queue.

**Pros:** Simple and fair - every process gets a turn, Good for interactive systems

**Cons:** Lots of context switching if quantum is too small. Too large → behaves like FCFS

### Priority Scheduling

**How it works:** Each process is assigned a priority

- Can be
  - Preemptive : higher priority interrupts
  - Non-preemptive : lower priority finishes first

**Pros:** Good for real time tasks or critical jobs

**Cons:** Starvation if low priority jobs never run, Aging is needed to prevent starvation

**Def : Aging** - Gradually increases the priority of a process the longer it waits.

**Def - Preemption** - OS's ability to interrupt/pause running process so that another can run

**CPU - I/O Bursts** : Pattern in which a process alternates between CPU bursts (doing computation) and I/O bursts (waiting for I/O operations like reading from disk or network)

- Most processes follow the cycle : short CPU bursts, long I/O bursts




Figure 5.1 Alternating sequence of CPU and I/O bursts.

**Turnaround Time** - The total time from when a process is submitted to when it completes.

⋮ Turnaround Time = Completion Time - Arrival Time

**Response Time** - The time from when a process is submitted to when it first starts execution

Response Time = First Run Time - Arrival Time

**Waiting Time** - The total time a process spends in the ready queue, waiting to get CPU time

Waiting Time = Turnaround Time - CPU Burst Time

**Throughput** - The number of processes completed per unit of time

Queue Type	Purpose	Common Algorithm	Priority Level
System	OS-level processes	Preemptive Priority	Very High
Interactive	Real-time apps (GUIs, shells)	Round Robin, SRTF	High
Interactive Edit	Editors, IDEs	Round Robin	Medium-High
Batch	Background tasks	FCFS, SJF	Medium-Low
Student / Low	Low-priority jobs	FCFS	Low

## Exponential Averaging Equation

In SJF, the scheduling must pick the process with the shortest next CPU burst. It predicts this time using the process's previous behavior.

$$\tau(n+1) = \alpha \cdot t(n) + (1 - \alpha) \cdot \tau(n)$$

- $\tau(n+1)$ : Predicted burst time for the next CPU burst

- $t(n)$ : Actual burst time of the current burst

- $\tau(n)$ : Predicted burst time for the current burst

- $\alpha$  (alpha): Smoothing factor between 0 and 1

This is a weighted average.  $\alpha$  controls how much weight is given to the recent burst time vs. the historical prediction

$a = 1 \rightarrow$  Use only the most recent burst → Reacts quickly but may be unstable

$a = 0 \rightarrow$  Use only the past predictions → Very stable, but doesn't learn

$0 < \alpha < 1 \rightarrow$  Weighted mix of past + present → Balance between adaptability and stability

## Preemption in Real Time

Real time schedulers are often preemptive so that critical tasks can interrupt less important ones to meet deadlines

- Hard real time systems always use preemptive scheduling
- Soft real time systems may use non-preemptive algorithms to simplify design

## Challenges in real time scheduling

- Priority inversion: A low priority process holds a resource needed by a high priority one
  - Solving using priority inheritance (temporarily raises the priority of a process to match the priority of another)
- Overhead from using too many preemptions
- Deadline overhead: When too many high priority tasks arrive at once, the system may drop or reject low priority jobs

## Real Time Scheduling

Two kinds of real time scheduling

- Hard Real Time

- Deadlines must be met, missing even one can cause system failure.

- Soft real time

- Deadlines are important, but missing them occasionally is okay

## Goals of Real Time Scheduling

- Determinism: Predictable execution timing

- Deadline guarantees: Meet task specific timing constraints

- Minimal jitter: Keeping timing consistent

- High priority for timecritical tasks

## Real Time Scheduling Algorithms

### 1. Rate Monotonic Scheduling (RMS)

The more frequently a task needs to run, the more important it is

- Priorities are fixed, they don't change during execution
- Works only if all tasks are
  - Periodic (runs on a fixed schedule)
  - Independent (no shared resources that cause blocking)
- Best for systems with predictable and repetitive jobs

### 2. Earliest Deadline First (EDF)

Always run the task that is closest to its deadline

- Priorities are dynamic - they change depending on the deadline

- Can handle both
  - Periodic tasks
  - Aperiodic tasks (come in randomly)

- Best for systems where timing is more flexible, or tasks arrive at random

## Short Term Scheduler

- Selects which ready process runs on the CPU next; runs very frequently; must be fast

## Long Term Scheduler

- Decided which jobs to enter the system from the job pool; runs less often; controls degree of multiprogramming

## Def - Preemptive Scheduling

- Happens when high priority process arrives or time quantum expires

## Def - Non Preemptive Scheduling

- Once process starts running it runs to completion or until it voluntarily yields

## Def - Priority Inversion

- low priority processes block high priority, one by holding a needed resource. Fixed with priority inheritance

## Def - Priority Inheritance

- low priority process, temporarily inherits the higher priority of a blocked process to prevent priority inversion, then reverts back after releasing the resource




Figure 5.8 Multilevel queue scheduling.